



AWS re:**Invent**

The AWS re:Invent logo is positioned on the left side of the slide. It features the word "AWS" in a black sans-serif font, followed by "re:" in a smaller orange font, and "Invent" in a large, bold black font. The background of the logo is a dense, abstract pattern of blue and orange hexagons and lines, resembling a circuit board or a network diagram.

DVO401

Deep Dive into Blue/Green Deployments on AWS

Andy Mui, Vlad Vlasceanu

October 2015

What to expect from the session

- Overview of common deployment risks
- Blue/green deployment concepts
- Benefits of blue/green with AWS
- Deploying apps using blue/green patterns on AWS
- Best practices for the data tier
- Cost optimization

Deployments are not easy

```
[stormtrooper1337@deathstar ~]$ cat deploy.sh
#!/bin/bash
# Date
#
# Last Updated: Luke S; Jan 13, 2010 ■ Introduction to Linux
# Desc: Use me to deploy code to app servers. I will install needed
# dependencies and ship the code. Before using me, please make sure
# your local unit tests pass. Backup directory is /backup/* in case
# something fails and you need to roll back. Talk to Tim if the host
# becomes unreachable. Good luck and feel free to improve me!
#
# Usage:
# deploy.sh <app version> <QA|STAGE|PROD>
#
APPDIR="/src/superlaser/$2"
HOSTNAME=$(hostname)
BEGIN="$(date)"
echo "Starting upgrade process - ${BEGIN}"
case $2 in
    QA)
        echo "Deploying QA code to ${HOSTNAME}"
;;
)
```

- Traditional environments favor in-place upgrades
- Resource constraints
- Downtime
- Dependencies
- Process coordination
- Difficult rollback

Common deployment risks

Challenges

- Application failure
- Infrastructure failure
- Capacity issues
- Scaling issues
- People failure
- Process failure
- Rollback issues

Business Impacts

- Downtime
- Data loss
- Bad customer experience
- Lost revenue
- Grumpy managers
- Burned out staff
- Wasted time/resources



Defining blue/green deployment on AWS

What is blue/green deployment?

“**Blue**”

(existing production environment)

“**Green**”

(parallel environment running a different version of the application)

“**Deployment**”

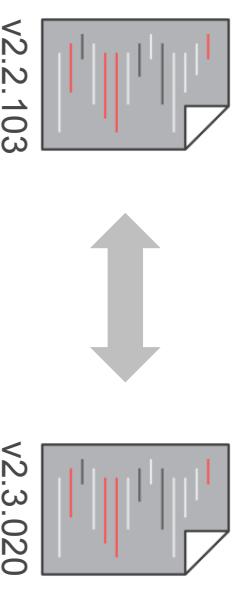
(ability to switch traffic between the two environments)

What is an environment?

Boundary for where things changed and what needs to be deployed

Examples:

DNS, load balancer

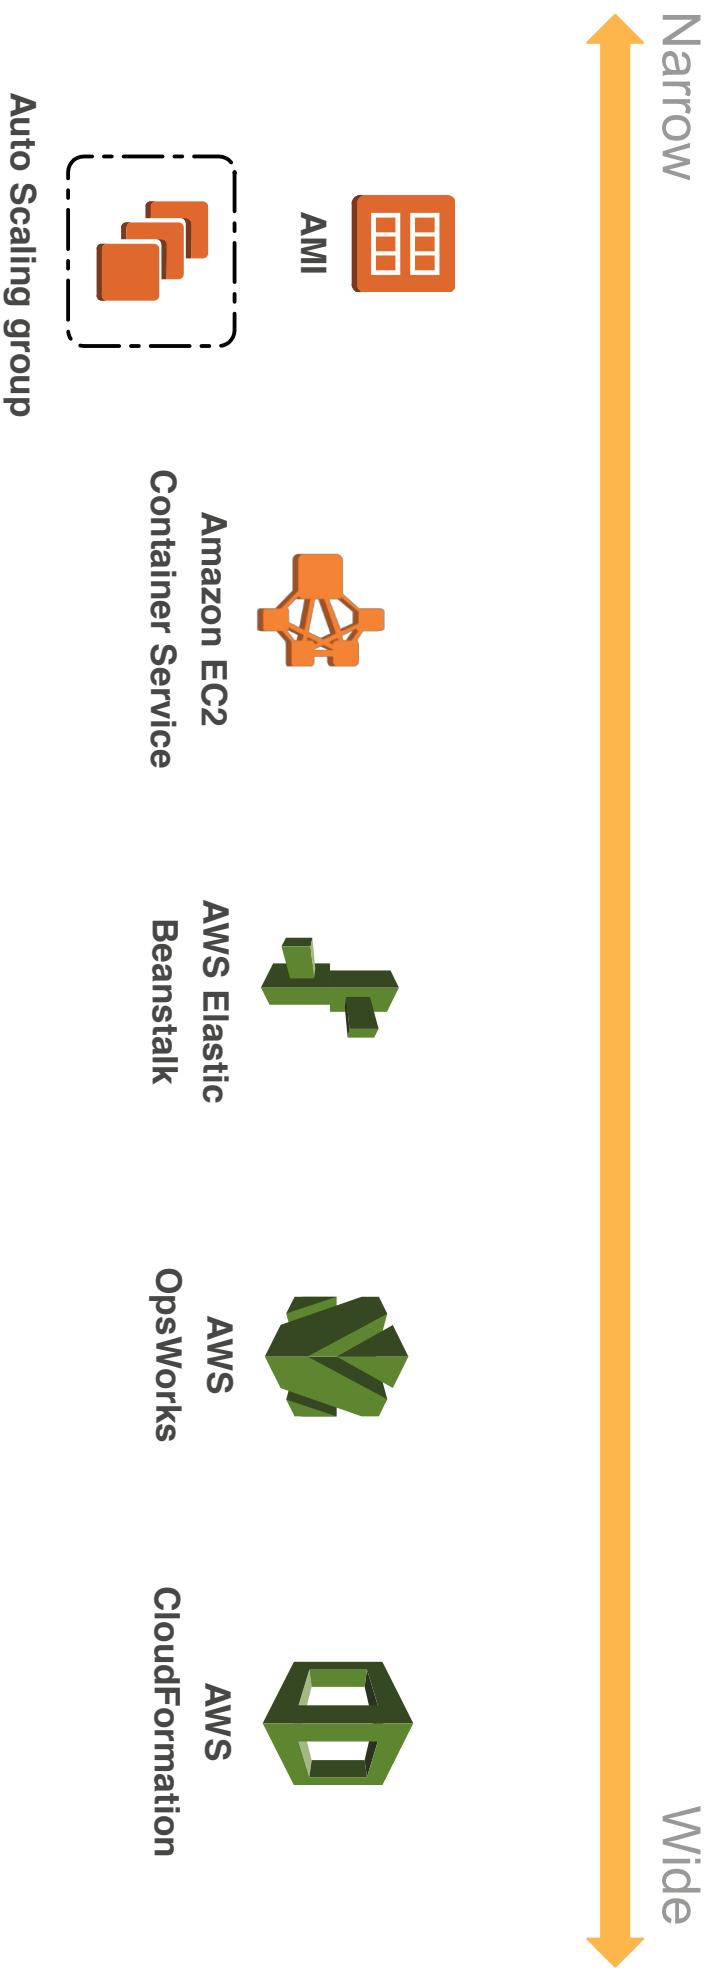


App component, app tier, microservice

v2.2.103

v2.3.020

Environment scope in AWS



Define your environment boundary

Factors	Criteria
Application architecture	Dependencies, loosely/tightly coupled
Organizational	Speed and number of iterations
Risk and complexity	Blast radius and impact of failed deployment
People	Expertise of teams
Process	Testing/QA, rollback capability
Cost	Operating budgets

Every deployment is different in scope and risk.

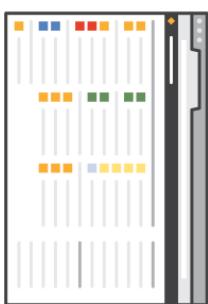
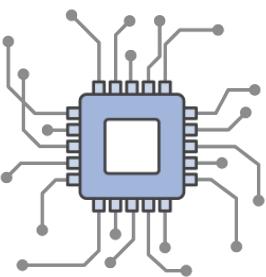
- Need a platform that is flexible and has powerful automation tools.



Benefits of blue/green deployment on AWS

AWS:

- Agile deployments
- Flexible options
- Scalable capacity
- Pay for what you use
- Automation capabilities



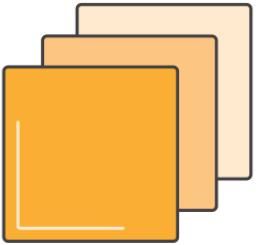


Deploying apps using blue/green patterns on AWS

Deploying apps using blue/green patterns

Using EC2 instances

1. Classic DNS cutover
2. Swap Auto Scaling groups
3. Swap launch configurations

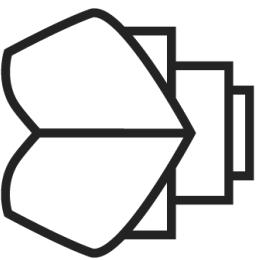


Using EC2 Container Service

1. Swap ECS services via DNS

2. Swap ECS services behind ELB

3. Swap ECS task definitions



Common thread: Environment automation

Deployment success depends on mitigating risk for:

- Application issues (functional)
- Application performance
- **People/process errors**
- **Infrastructure failure**
- Rollback capability
- Large costs

Strength of automation platform

CloudFormation most comprehensive automation platform

- Scope stacks from network to software
- Control higher-level automation services:

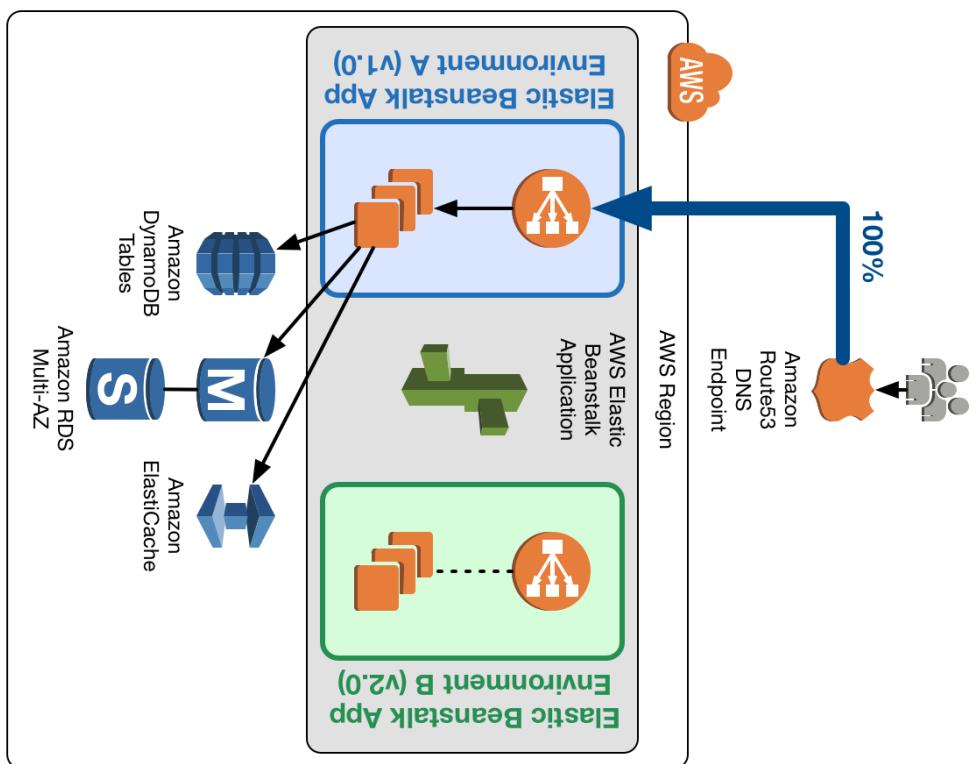
Elastic Beanstalk, ECS, OpsWorks, Auto Scaling

Blue/green deployment patterns address these risks differently

Patterns: classic DNS cutover

Deployment process:

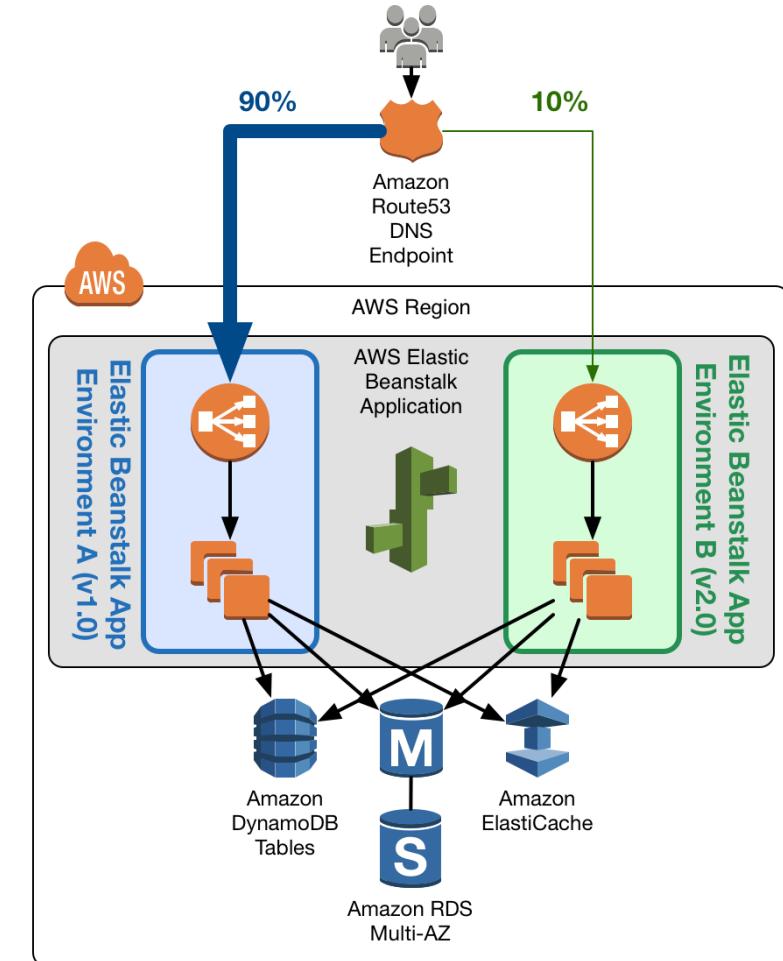
- Start with **current** app environment
- Deploy a **new** app environment
- Test the green stack
- Gradually cut traffic over via DNS
- Monitor your environments
- If needed, roll back to blue



Patterns: classic DNS cutover

Deployment process:

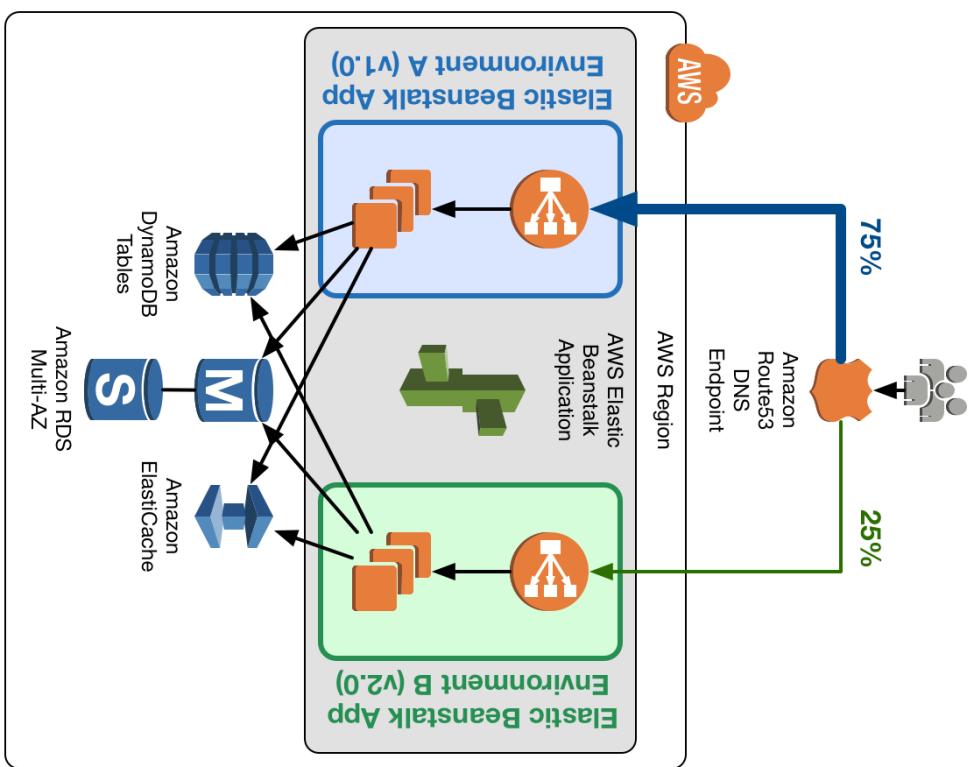
- Start with **current** app environment
- Deploy a **new** app environment
- **Test** the green stack
- Gradually **cut traffic over** via DNS
- **Monitor** your environments
- If needed, **roll back** to blue



Patterns: classic DNS cutover

Deployment process:

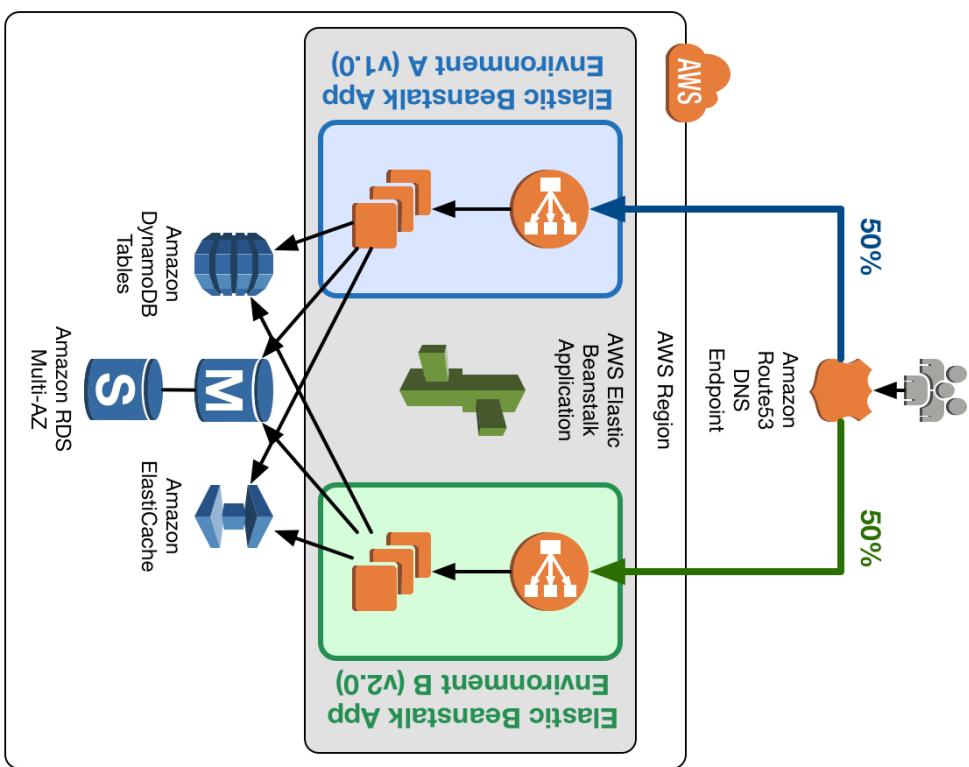
- Start with **current** app environment
- Deploy a **new** app environment
- **Test** the green stack
- Gradually cut traffic over via DNS
- Monitor your environments
- If needed, roll back to blue



Patterns: classic DNS cutover

Deployment process:

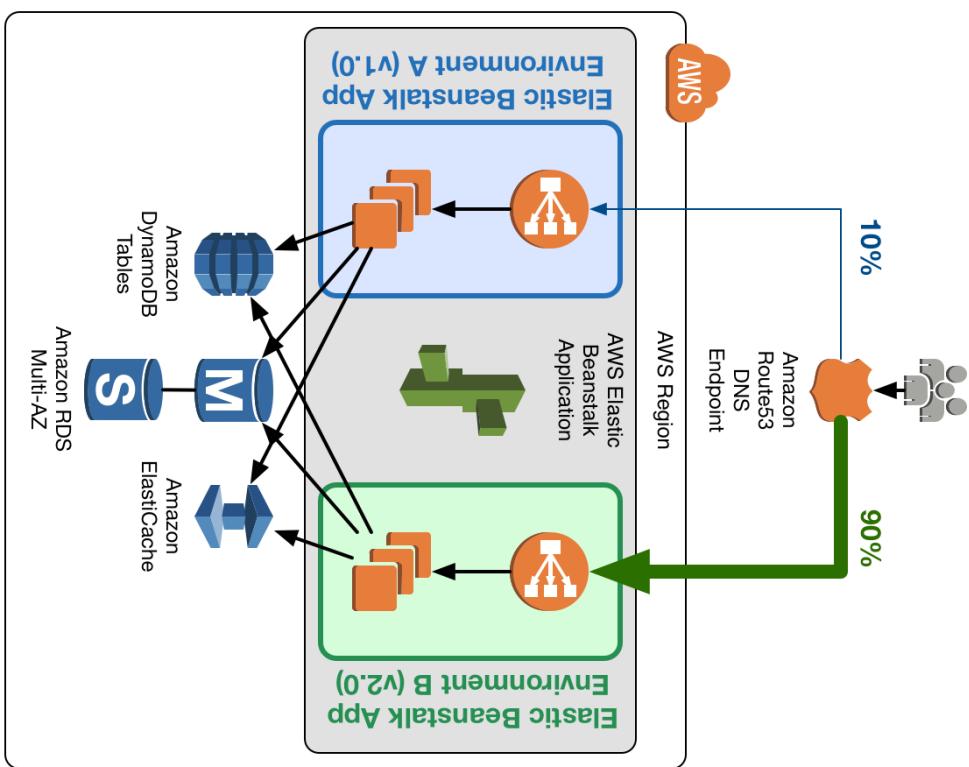
- Start with **current** app environment
- Deploy a **new** app environment
- **Test** the green stack
- Gradually cut traffic over via DNS
- **Monitor** your environments
- If needed, roll back to blue



Patterns: classic DNS cutover

Deployment process:

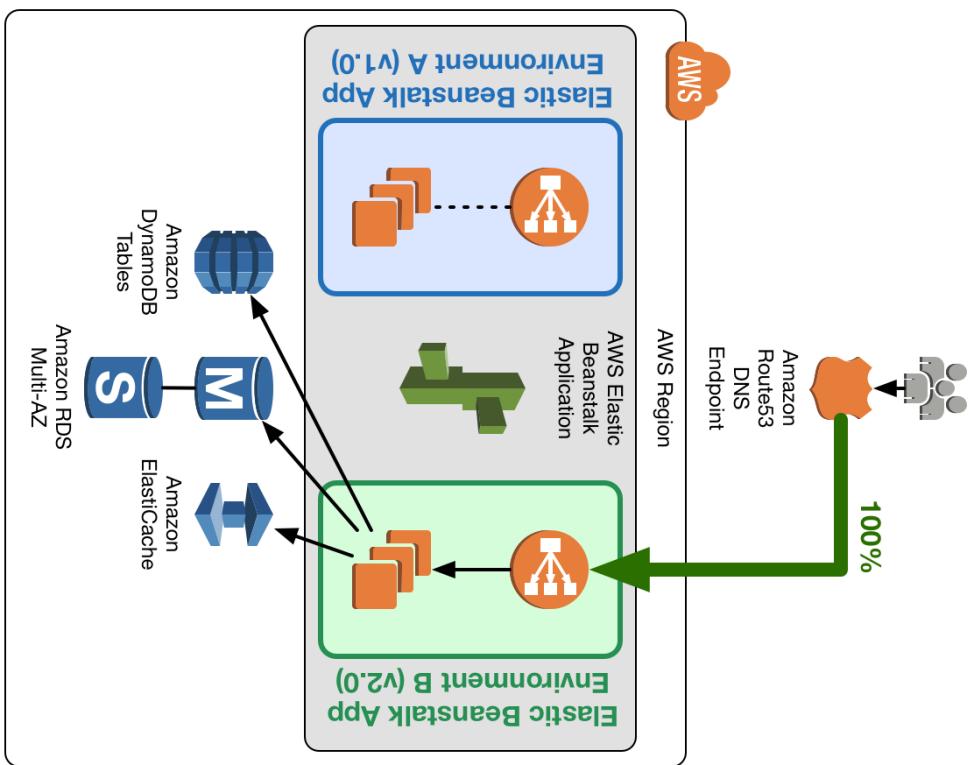
- Start with **current** app environment
- Deploy a **new** app environment
- **Test** the green stack
- Gradually cut traffic over via DNS
- **Monitor** your environments
- If needed, roll back to blue



Patterns: classic DNS cutover

Deployment process:

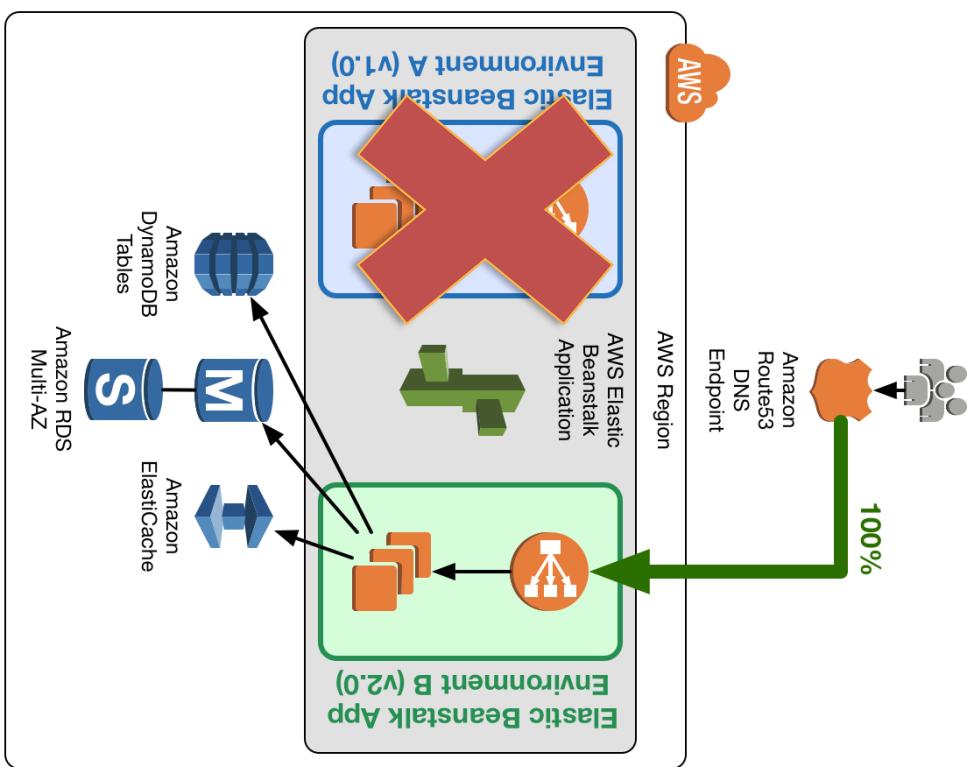
- Start with **current** app environment
- Deploy a **new** app environment
- **Test** the green stack
- Gradually cut traffic over via DNS
- **Monitor** your environments
- If needed, roll back to blue



Patterns: classic DNS cutover

Deployment process:

- Start with **current** app environment
- Deploy a **new** app environment
- **Test** the green stack
- Gradually cut traffic over via DNS
- **Monitor** your environments
- If needed, roll back to blue



Automating your environment

- Use **CloudFormation** templates to model your environment
 - Version-control your templates
 - Use Elastic Beanstalk or OpsWorks to model your applications inside the template
 - Update CloudFormation stack from updated template containing green environment
- ```
"Resources": {
 "myApp": { "Type": "AWS::ElasticBeanstalk::Application" },
 "myConfigTemplate": {
 "Type": "AWS::ElasticBeanstalk::ConfigurationTemplate"
 },
 "myBlueAppVersion": {
 "Type": "AWS::ElasticBeanstalk::Applicationversion"
 },
 "myBlueEnvironment": {
 "Type": "AWS::ElasticBeanstalk::Environment"
 },
 "myBlueEndpoint": { "Type": "AWS::Route53::Recordset" },
 "myGreenAppVersion": {
 "Type": "AWS::ElasticBeanstalk::Applicationversion"
 },
 "myGreenEnvironment": {
 "Type": "AWS::ElasticBeanstalk::Environment"
 },
 "myGreenEndpoint": { "Type": "AWS::Route53::Recordset" }
},
...
```

# Amazon Route 53 weighted DNS switching

- AWS Elastic Beanstalk environment endpoint swap
- DNS record time-to-live (TTL)
  - Reaction time = (TTL × no. of DNS caches) + Route53 propagation time, up to 1min
  - Beware of misbehaving DNS clients
- Auto Scaling and Amazon Elastic Load Balancing (ELB) need time to scale
- Measurable metrics
  - ELB: Latency, SurgeQueueLength, SpillOverCount, BackendConnectionErrors
  - Your application metrics
- Your deployment goals

# Amazon Route 53 weighted DNS switching

## Using CloudFormation:

- Update template record sets with initial weighting information
- Consider using parameters for the weight values – reuse the same template
- Update CloudFormation stack with new weighting

```
"myBlueEndpoint": {
 "Type": "AWS::Route53::RecordSet",
 "Properties": {
 "HostedZoneId": { "Ref": "parameterHostedZoneId" },
 "Name": "www.example.com.", "Type": "CNAME", "TTL": "60",
 "SetIdentifier": "stack-blue", "Weight": "90",
 "ResourceRecords": [
 { "Fn::GetAtt": ["myBlueEnvironment", "EndpointURL"] }
] } },
 "myGreenEndpoint": {
 "Type": "AWS::Route53::RecordSet",
 "Properties": {
 "HostedZoneId": { "Ref": "parameterHostedZoneId" },
 "Name": "www.example.com.", "Type": "CNAME", "TTL": "60",
 "SetIdentifier": "stack-green", "Weight": "10",
 "ResourceRecords": [
 { "Fn::GetAtt": ["myGreenEnvironment", "EndpointURL"] }
] } }
```



# Pattern review: Classic DNS cutover

| Risk category           | Mitigation level | Reasoning                                       |
|-------------------------|------------------|-------------------------------------------------|
| Application issues      | Great            | Facilitates canary analysis                     |
| Application performance | Great            | Gradual switch, traffic split management        |
| People/process errors   | Good             | Depends on automation framework                 |
| Infrastructure failure  | Good             | Depends on automation framework                 |
| Rollback                | Fair             | DNS TTL complexities (reaction time, flip/flop) |
| Cost                    | Great            | Optimized via auto-scaling                      |

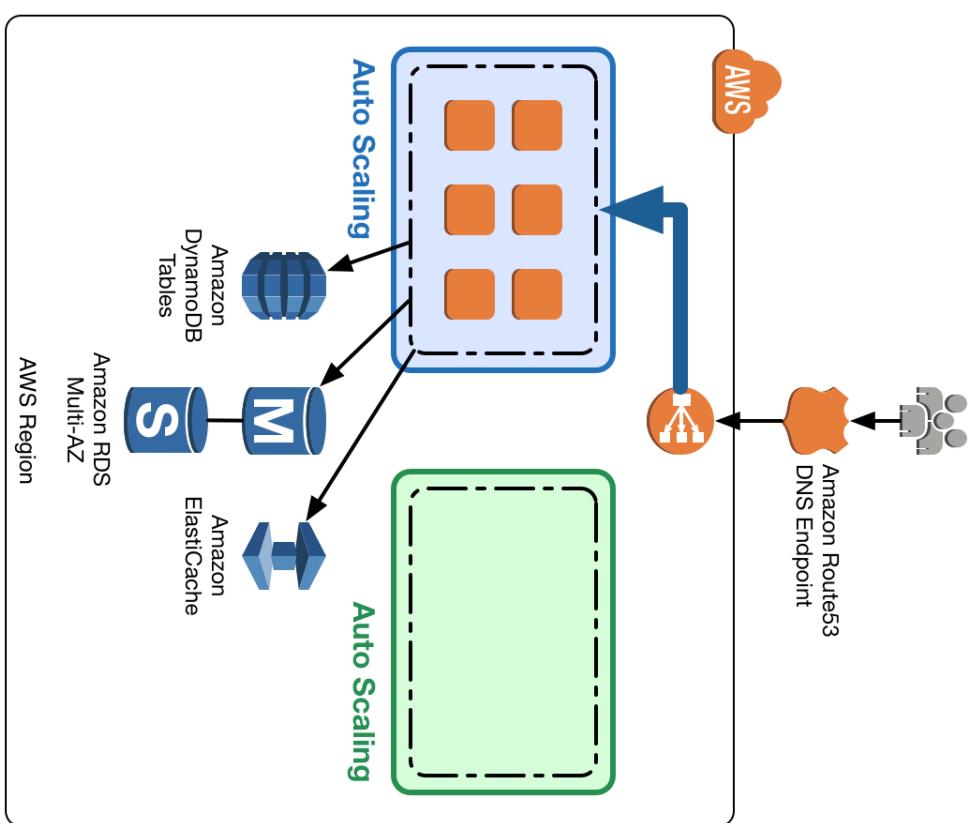


**Let's remove DNS from the  
solution...**

# Pattern: swap Auto Scaling Groups

## Deployment process:

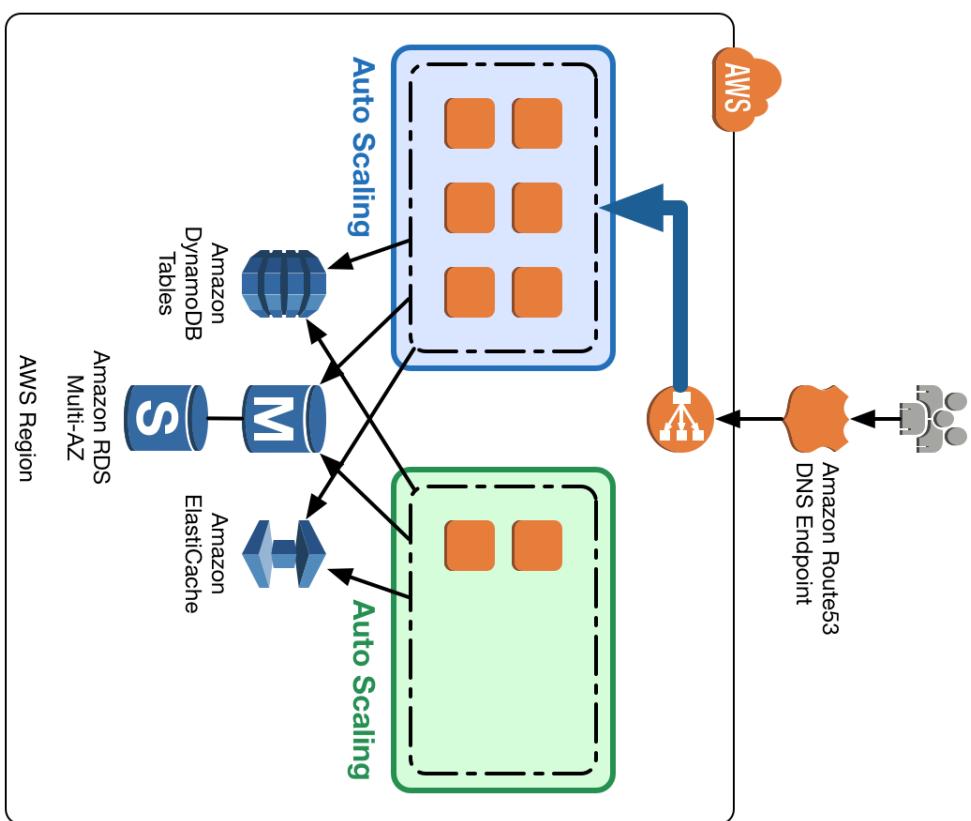
- Amazon Elastic Load Balancer (ELB) outside the environment boundary
- Start with **current** Auto Scaling Group (ASG)
- Deploy & scale out **new** ASG
- Test green stack
- Register green ASG with ELB
- Remove blue ASG from ELB



# Pattern: swap Auto Scaling Groups

## Deployment process:

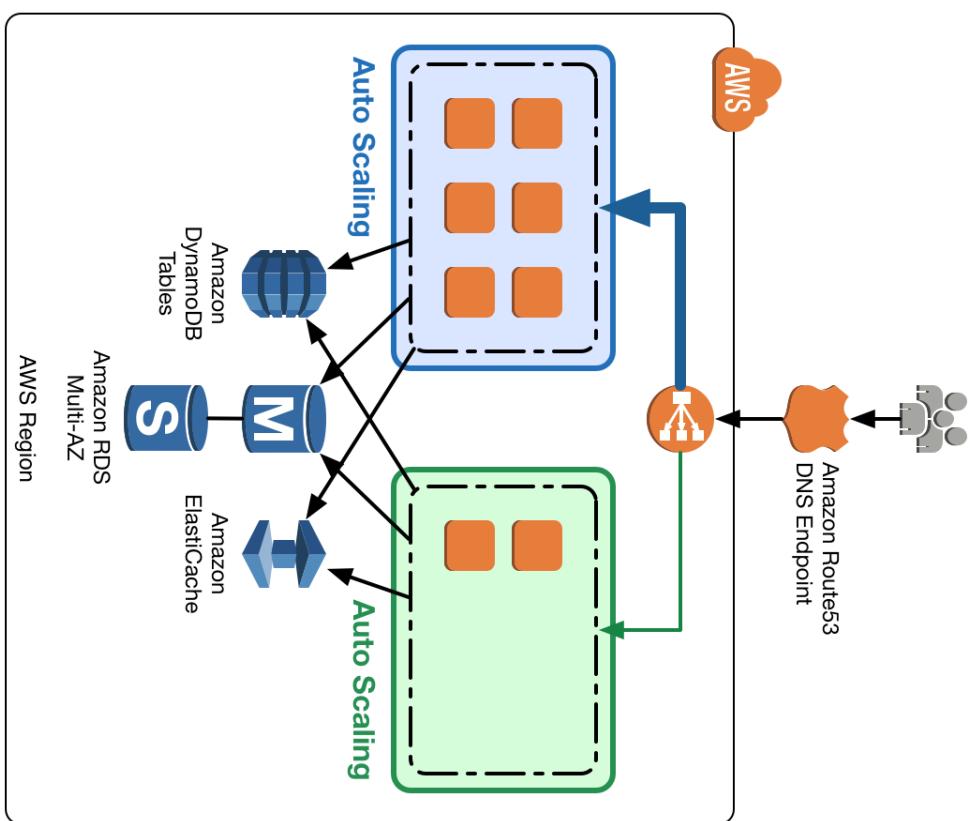
- Amazon Elastic Load Balancer (ELB) outside the environment boundary
- Start with **current** Auto Scaling Group (ASG)
- Deploy & scale out **new** ASG
- Test green stack
- Register green ASG with ELB
- Remove blue ASG from ELB



# Pattern: swap Auto Scaling Groups

## Deployment process:

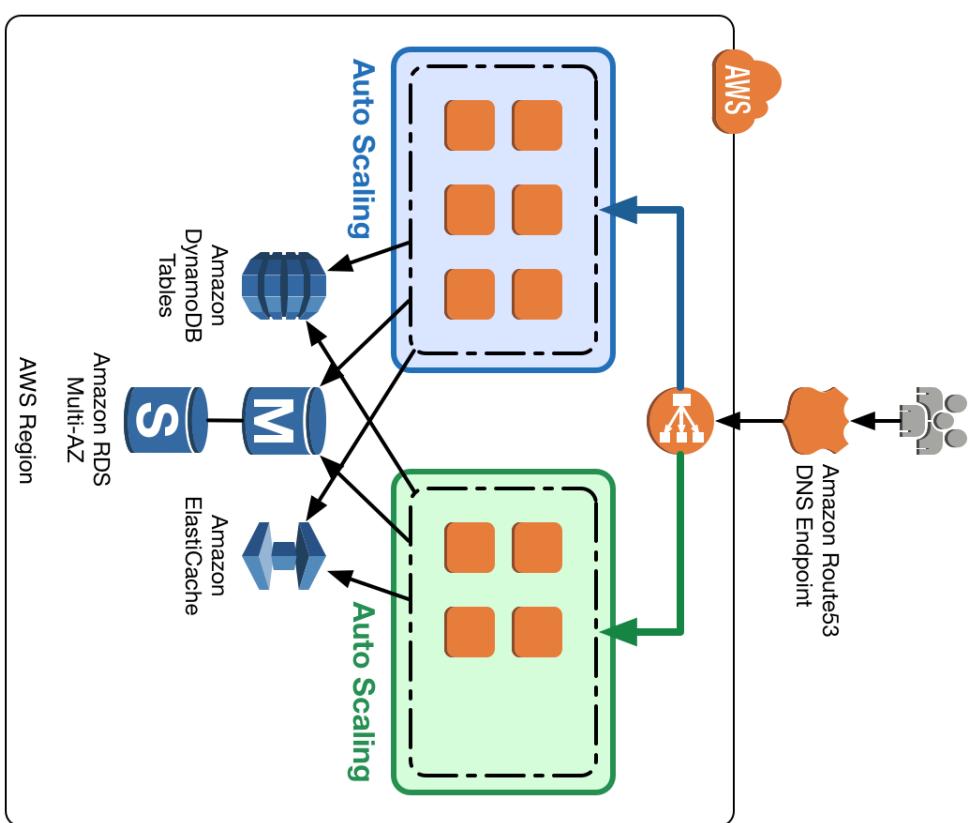
- Amazon Elastic Load Balancer (ELB) outside the environment boundary
- Start with **current** Auto Scaling Group (ASG)
- Deploy & scale out **new** ASG
- Test green stack
- Register green ASG with ELB
- Remove blue ASG from ELB



# Pattern: swap Auto Scaling Groups

## Deployment process:

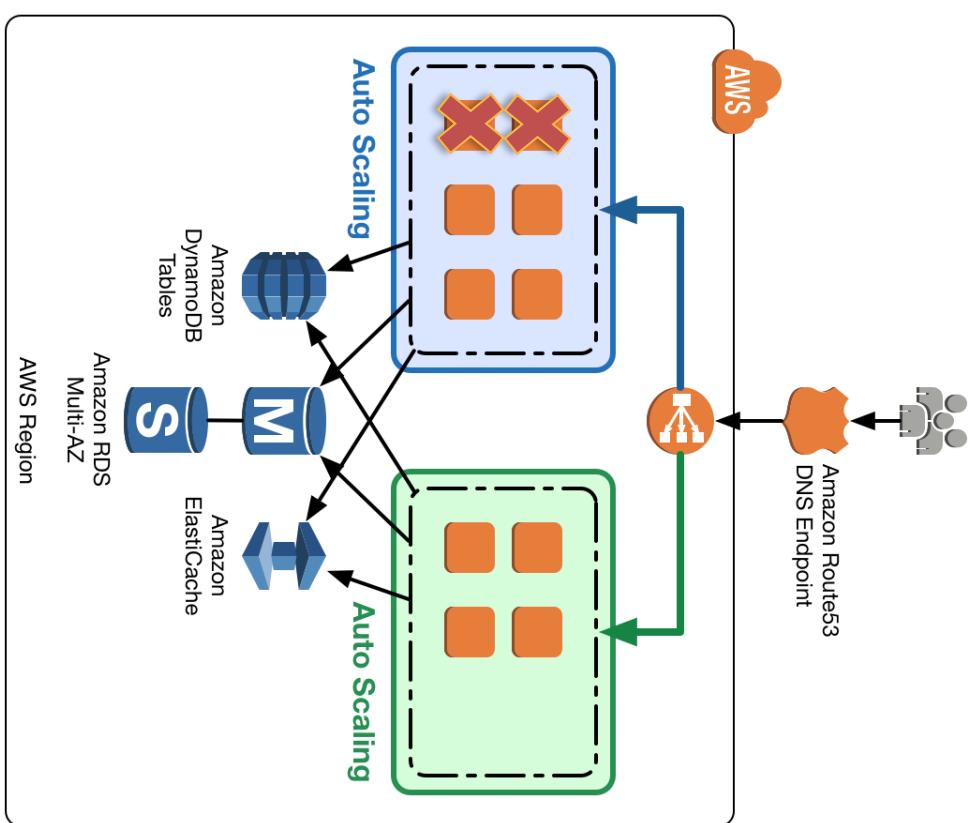
- Amazon Elastic Load Balancer (ELB) outside the environment boundary
- Start with **current** Auto Scaling Group (ASG)
- Deploy & scale out **new** ASG
- Test green stack
- Register green ASG with ELB
- Remove blue ASG from ELB



# Pattern: swap Auto Scaling Groups

## Deployment process:

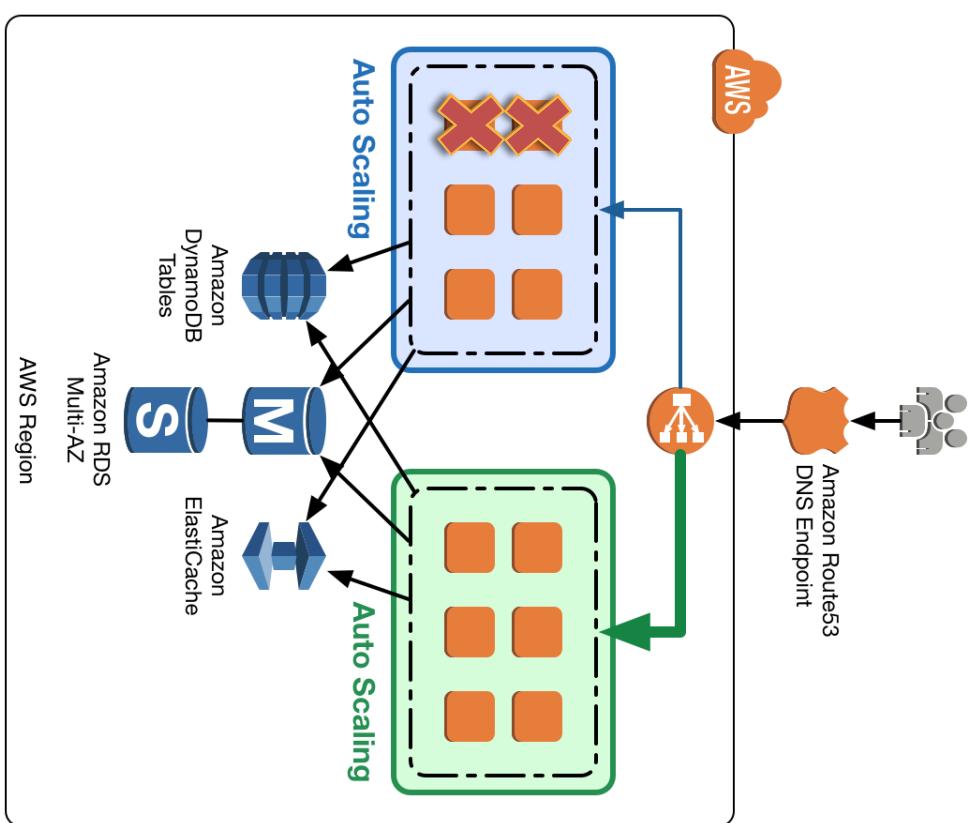
- Amazon Elastic Load Balancer (ELB) outside the environment boundary
- Start with **current** Auto Scaling Group (ASG)
- Deploy & scale out **new** ASG
- Test green stack
- Register green ASG with ELB
- Remove blue ASG from ELB



# Pattern: swap Auto Scaling Groups

## Deployment process:

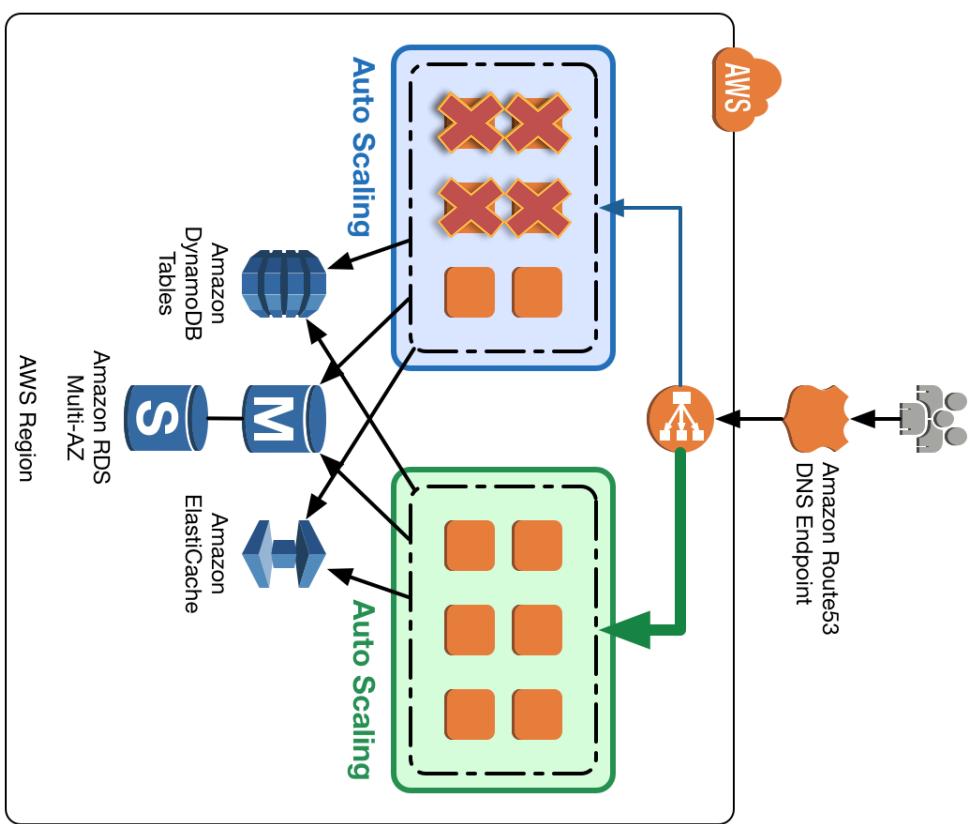
- Amazon Elastic Load Balancer (ELB) outside the environment boundary
- Start with **current** Auto Scaling Group (ASG)
- Deploy & scale out **new** ASG
- Test green stack
- Register green ASG with ELB
- Remove blue ASG from ELB



# Pattern: swap Auto Scaling Groups

## Deployment process:

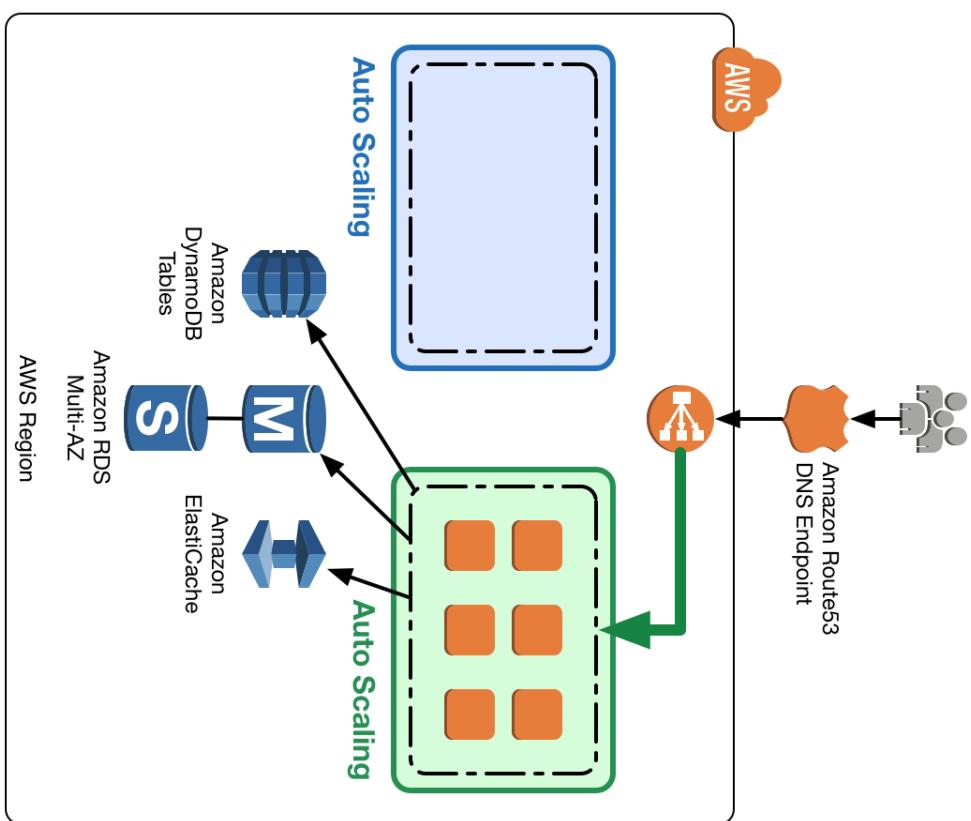
- Amazon Elastic Load Balancer (ELB) outside the environment boundary
- Start with **current** Auto Scaling Group (ASG)
- Deploy & scale out **new** ASG
- Test green stack
- Register green ASG with ELB
- Remove blue ASG from ELB



# Pattern: swap Auto Scaling Groups

## Deployment process:

- Amazon Elastic Load Balancer (ELB) outside the environment boundary
- Start with **current** Auto Scaling Group (ASG)
- Deploy & scale out **new** ASG
- Test green stack
- Register green ASG with ELB
- Remove blue ASG from ELB



# Swapping Auto Scaling groups behind ELB

- Register with ELB:
  - One or more EC2 instances
  - One or more Auto Scaling groups
- Least outstanding requests algorithm favors green ASG instances for new connections
  - \$ aws autoscaling set-desired-capacity \  
--auto-scaling-group-name "green-asg" \  
--desired-capacity X
- Connection draining - gracefully stop receiving traffic
  - \$ aws autoscaling detach-load-balancers \  
--auto-scaling-group-name "blue-asg" \  
--load-balancer-names "my-app-elb"
- Scale out green ASG before ELB registration
  - \$ aws autoscaling enter-standby \  
--instance-ids i-xxxxxxxxx \  
--auto-scaling-group-name "blue-asg" \  
--should-decrement-desired-capacity
- Put blue instances in standby

# Pattern review: Swap Auto Scaling groups

| Risk category           | Mitigation level | Reasoning                                                   |
|-------------------------|------------------|-------------------------------------------------------------|
| Application issues      | Great            | Facilitates canary analysis w/ additional ELB               |
| Application performance | Good             | Traffic split management, but less granular, pre-warmed ELB |
| People/process errors   | Good             | Depends on automation framework                             |
| Infrastructure failure  | Great            | Auto-scaling                                                |
| Rollback                | Great            | No DNS complexities                                         |
| Cost                    | Great            | Optimized via auto-scaling                                  |

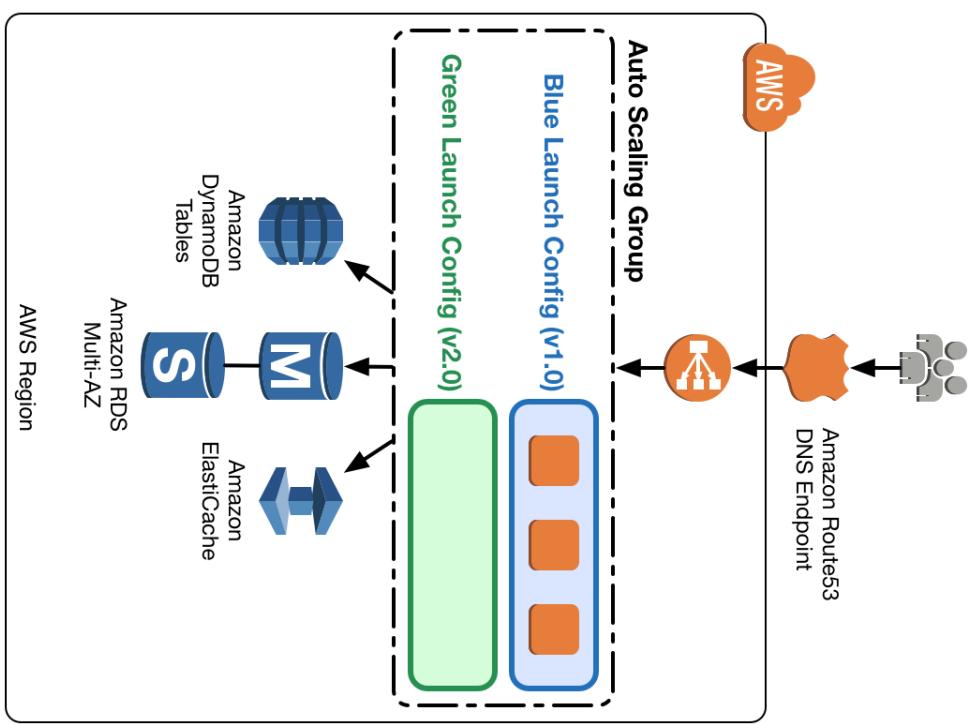


**Let's reduce the environment  
boundary further...**

# Pattern: swap Launch Configurations

## Deployment process:

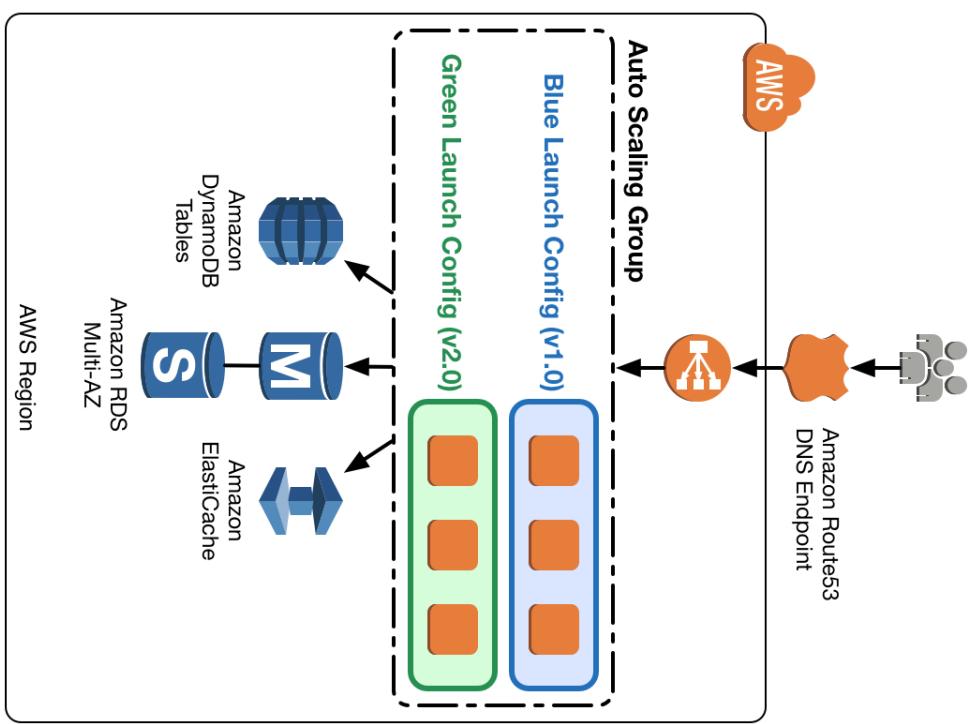
- Start with **current** ASG & Launch Configuration behind the ELB
- **Attach updated green** Launch Configuration to the ASG
- Grow the ASG gradually to 2x original size
- Shrink the ASG back to original size
- For more control, put old instances into Standby



# Pattern: swap Launch Configurations

## Deployment process:

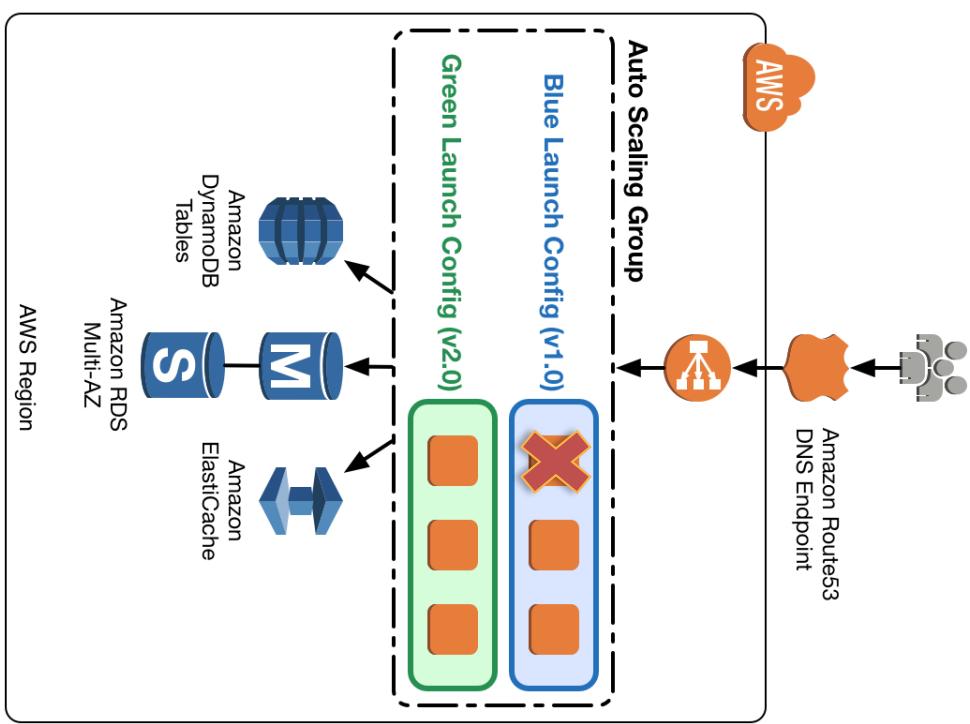
- Start with **current** ASG & Launch Configuration behind the ELB
- **Attach updated green** Launch Configuration to the ASG
- Grow the ASG gradually to 2x original size
- Shrink the ASG back to original size
- For more control, put old instances into Standby



# Pattern: swap Launch Configurations

## Deployment process:

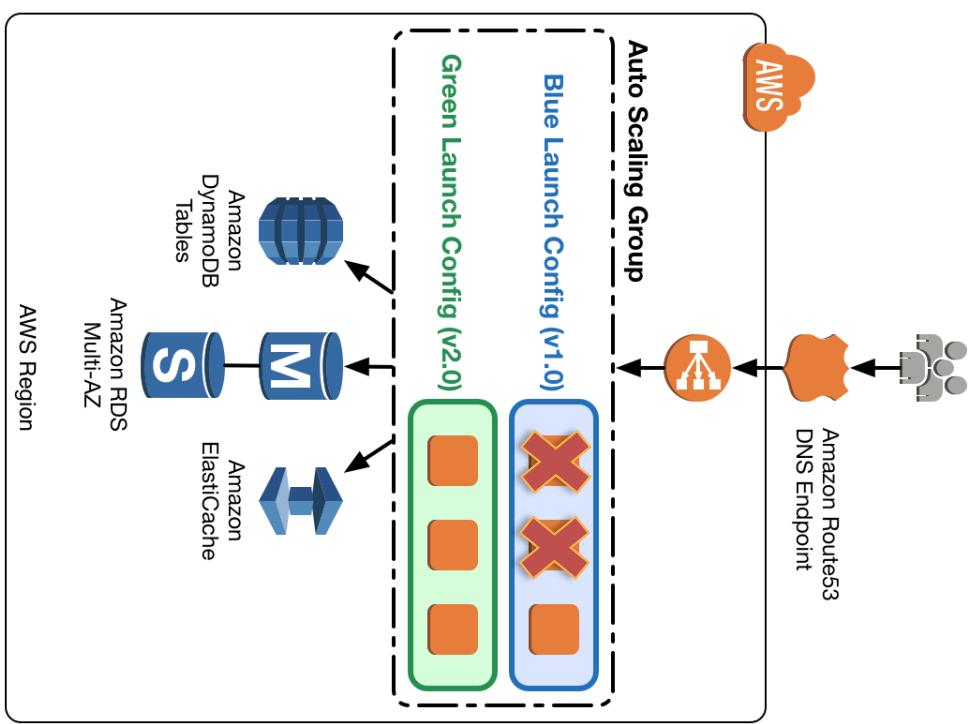
- Start with **current** ASG & Launch Configuration behind the ELB
- **Attach updated green** Launch Configuration to the ASG
- Grow the ASG gradually to 2x original size
- Shrink the ASG back to original size
- For more control, put old instances into Standby



# Pattern: swap Launch Configurations

## Deployment process:

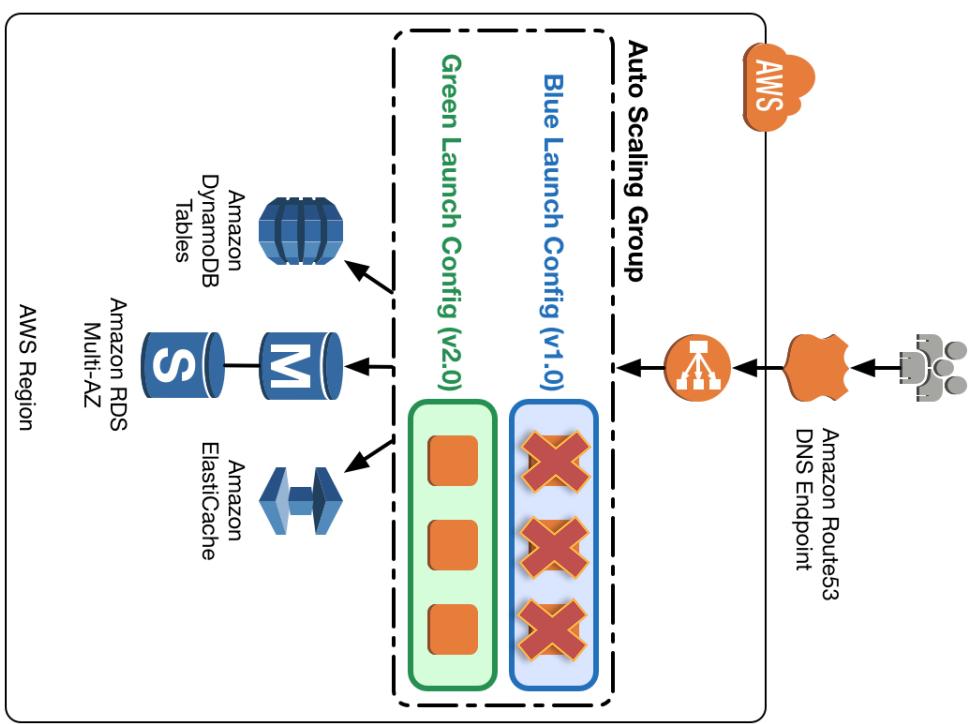
- Start with **current** ASG & Launch Configuration behind the ELB
- **Attach updated green** Launch Configuration to the ASG
- Grow the ASG gradually to 2x original size
- Shrink the ASG back to original size
- For more control, put old instances into Standby



# Pattern: swap Launch Configurations

## Deployment process:

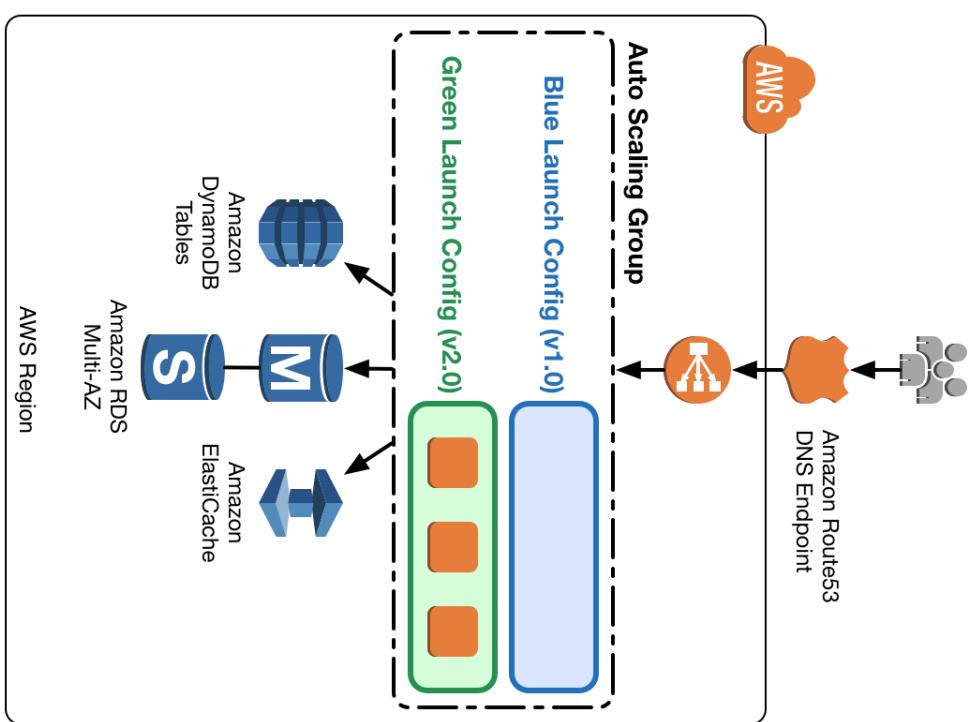
- Start with **current** ASG & Launch Configuration behind the ELB
- **Attach** updated **green** Launch Configuration to the ASG
- Grow the ASG gradually to 2x original size
- Shrink the ASG back to original size
- For more control, put old instances into Standby



# Pattern: swap Launch Configurations

## Deployment process:

- Start with **current** ASG & Launch Configuration behind the ELB
- **Attach updated green** Launch Configuration to the ASG
- Grow the ASG gradually to 2x original size
- Shrink the ASG back to original size
- For more control, put old instances into Standby



# Swapping launch configurations

- **Launch configurations:**  
Blueprints for ASG instance provisioning, each ASG points to exactly one
- **Scale-out & replacement:**  
Events will use the attached (**green**) launch configuration to provision instances
- **Scale-in:**  
ASG scale-in events will terminate instances with oldest launch configuration first **while** trying to keep capacity in AZs balanced
- May need to address AZ imbalances separately
- **Temporarily remove instances from ASG**  
Place specific ASG instances (**blue**) into standby – stop receiving traffic

# Pattern review: swap Launch Configurations

| Risk Category           | Mitigation Level | Reasoning                                                        |
|-------------------------|------------------|------------------------------------------------------------------|
| Application Issues      | Fair             | Detection of errors/issues in a heterogeneous fleet is complex   |
| Application Performance | Fair             | Less granular traffic split, initial traffic load                |
| People/Process Errors   | Good             | Depends on automation framework                                  |
| Infrastructure Failure  | Great            | Auto-Scaling                                                     |
| Rollback                | Great            | No DNS complexities                                              |
| Cost                    | Good             | Optimized via auto-scaling, but initial scale-out overprovisions |



What if we're running  
containers?

# EC2 Container Service

**Environment boundary:**

Docker container -> Task definition

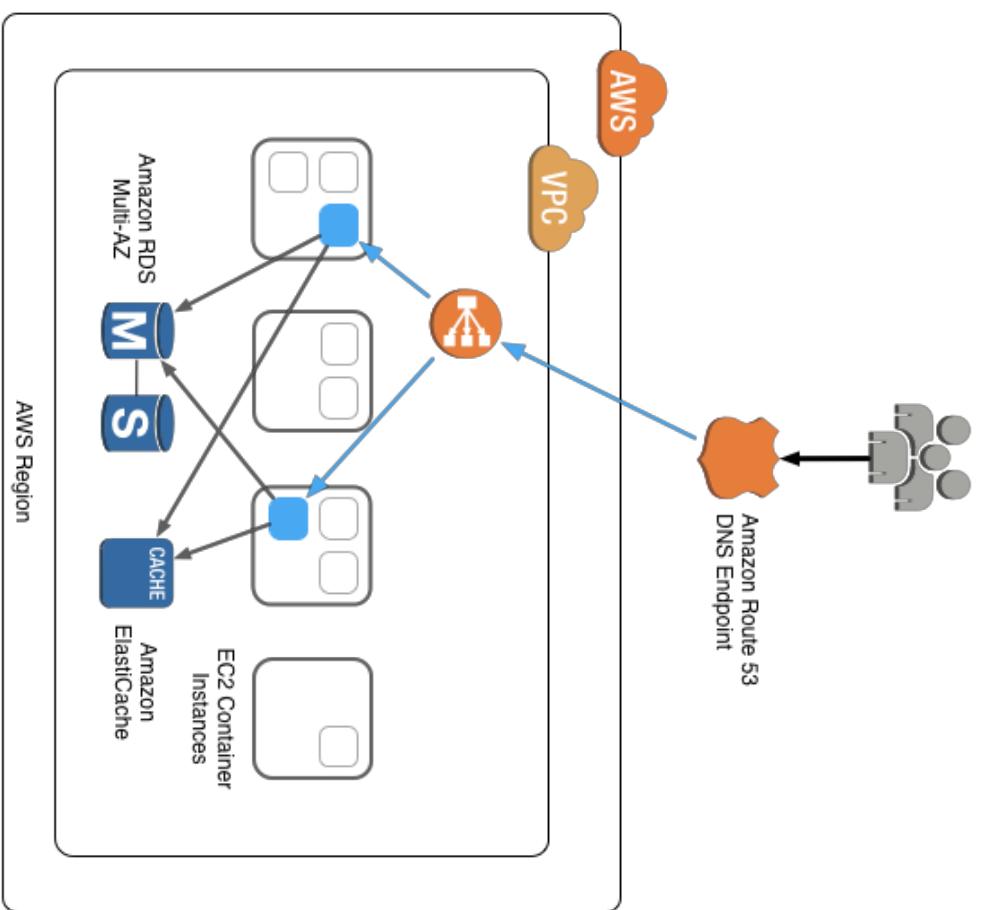
**Deployment options:**

1. Blue/green ECS services, DNS update
2. Blue/green ECS services, shared ELB
3. ECS update

# Pattern: Swap ECS services with DNS

## Deployment process:

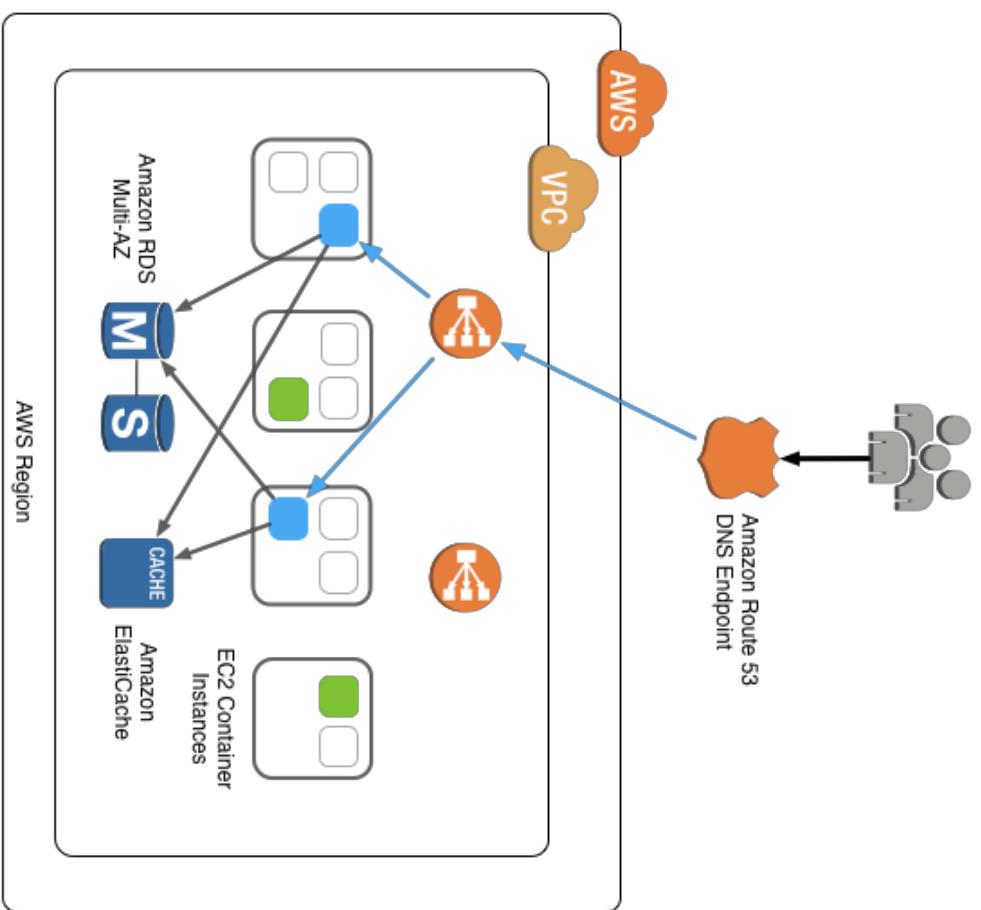
- Start with **blue** service composed of a task definition and ELB
- Create new task definition based on new version of Docker image and a new ELB
- Create **green** service with new task definition and ELB
- Update Route 53 alias record to direct traffic to new ELB endpoint
- Clean up blue service resources when no longer needed



# Pattern: Swap ECS services with DNS

## Deployment process:

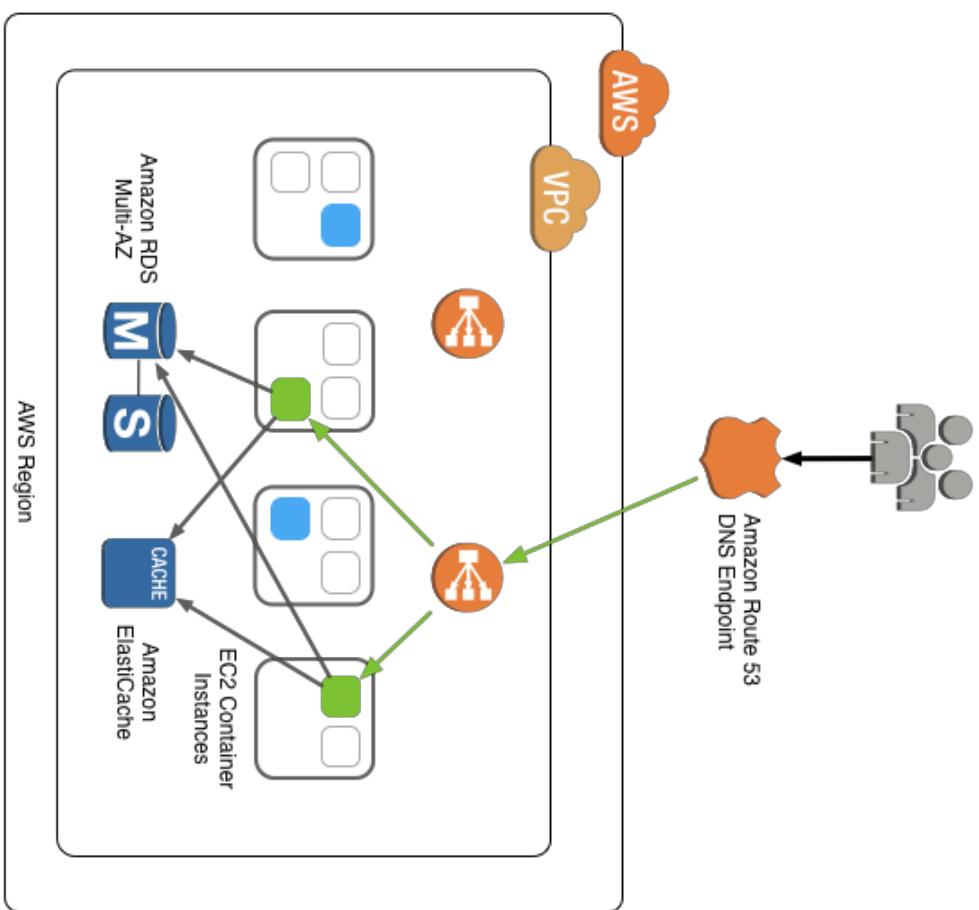
- Start with **blue** service composed of a task definition and ELB
- Create new task definition based on new version of Docker image and a new ELB
- Create **green** service with new task definition and ELB
- Update Route 53 alias record to direct traffic to new ELB endpoint
- Clean up blue service resources when no longer needed



# Pattern: Swap ECS services with DNS

## Deployment process:

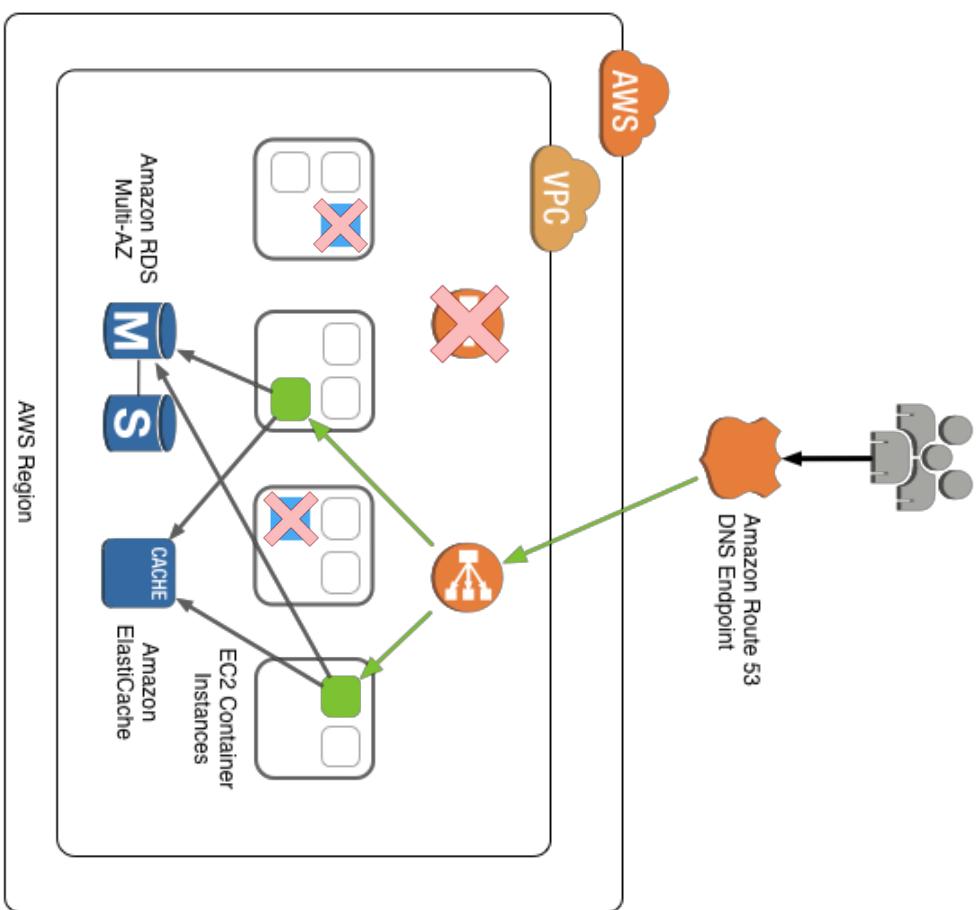
- Start with **blue** service composed of a task definition and ELB
  - Create new task definition based on new version of Docker image and a new ELB
  - Create **green** service with new task definition and ELB
  - Update Route 53 alias record to direct traffic to new ELB endpoint
  - Clean up blue Service resources when no longer needed



# Pattern: Swap ECS services with DNS

## Deployment process:

- Start with **blue** service composed of a task definition and ELB
- Create new task definition based on new version of Docker image and a new ELB
- Create **green** service with new task definition and ELB
- Update Route 53 alias record to direct traffic to new ELB endpoint
- Clean up blue service resources when no longer needed



# Swap ECS services with DNS

web.json

## 1. Create task definition

```
$ aws ecs register-task-definition \
--cli-input-json file:///Code/web.json
```

## 2. Create service and ELB

```
$ aws ecs create-service \
--cluster dice-demo \
--service-name web-service-v2 \
--task-definition web-tier \
--load-balancers LoadBalancerName=web-tier-
v2,containerName=web-tier,containerPort=8000 \
--desired-count 2 \
--role "ecsserviceRole"
```

# Swap ECS services with DNS

dns.json

## 3. Route traffic to **green** ELB endpoint

```
$ aws route53 change-resource-record-sets \
--hosted-zone-id zxxxxxxxxxxxxx \
--change-batch file://Code/dns.json
```

```
{
 "Comment": "Update alias record to route
traffic from Blue to Green ELB for Dice Demo
ECS Service",
 "Changes": [
 {"Action": "UPSERT",
 "ResourceRecordSet": {
 "Name": "demo.example.com",
 "Type": "A",
 "AliasTarget": {
 "HostedZoneId": "ZXXXXXXXXXXXXX",
 "DNSName": "dice-demo-green-
1234567890.us-west-2.elb.amazonaws.com",
 "EvaluateTargetHealth": false
 }
 }
]
}
```

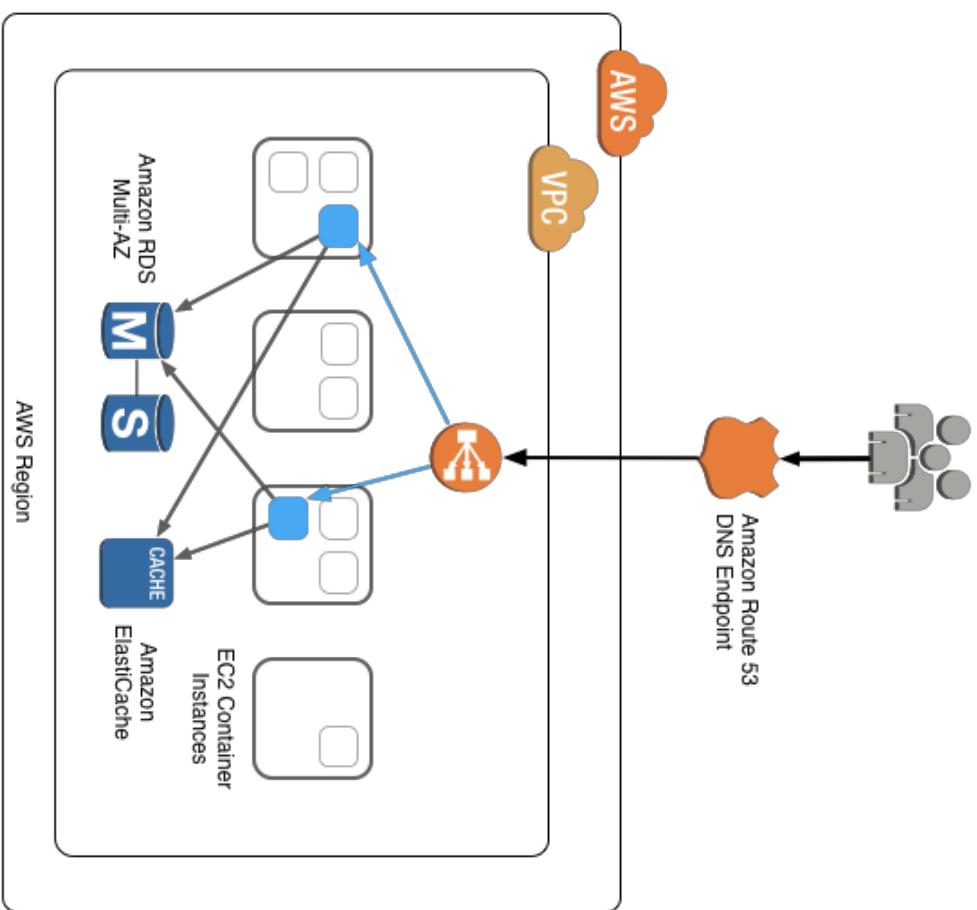
# Pattern review: Swap ECS services with DNS

| Risk category           | Mitigation level | Reasoning                                                        |
|-------------------------|------------------|------------------------------------------------------------------|
| Application issues      | Great            | Full cutover or weighted canary analysis                         |
| Application performance | Good             | ELB may require prewarm, CloudWatch, and automation to scale ECS |
| People/process errors   | Great            | Simple process                                                   |
| Infrastructure failure  | Good             | Leverage CloudWatch, Auto Scaling, ELB                           |
| Rollback                | Fair             | DNS TTL complexities (reaction time, flip/flop)                  |
| Cost                    | Fair             | Require enough cluster resources to accommodate new service      |

# Pattern: Swap ECS services with ELB

## Deployment process:

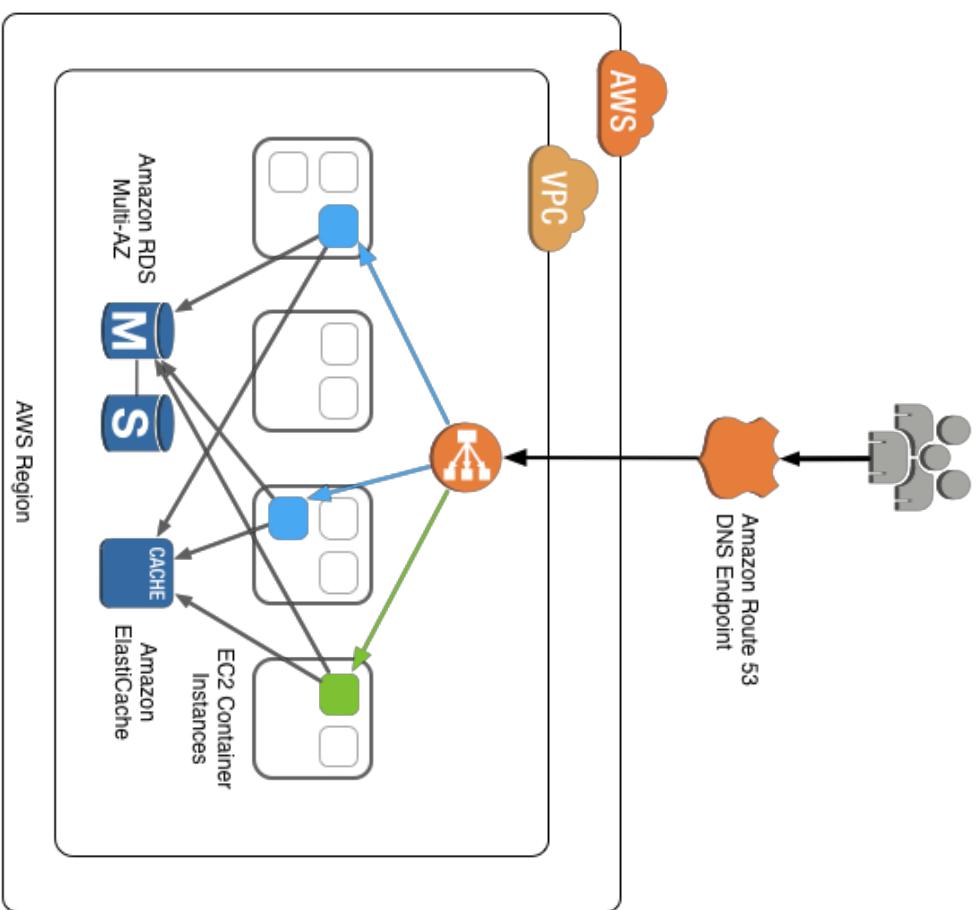
- Start with **blue** service composed of a task definition and ELB
- Create new task definition based on new version of Docker image
- Create **green** service with new task definition and map to existing ELB
- Scale up green service by incrementing number of tasks
- Decommission blue service by setting task count to 0



# Pattern: Swap ECS services with ELB

## Deployment process:

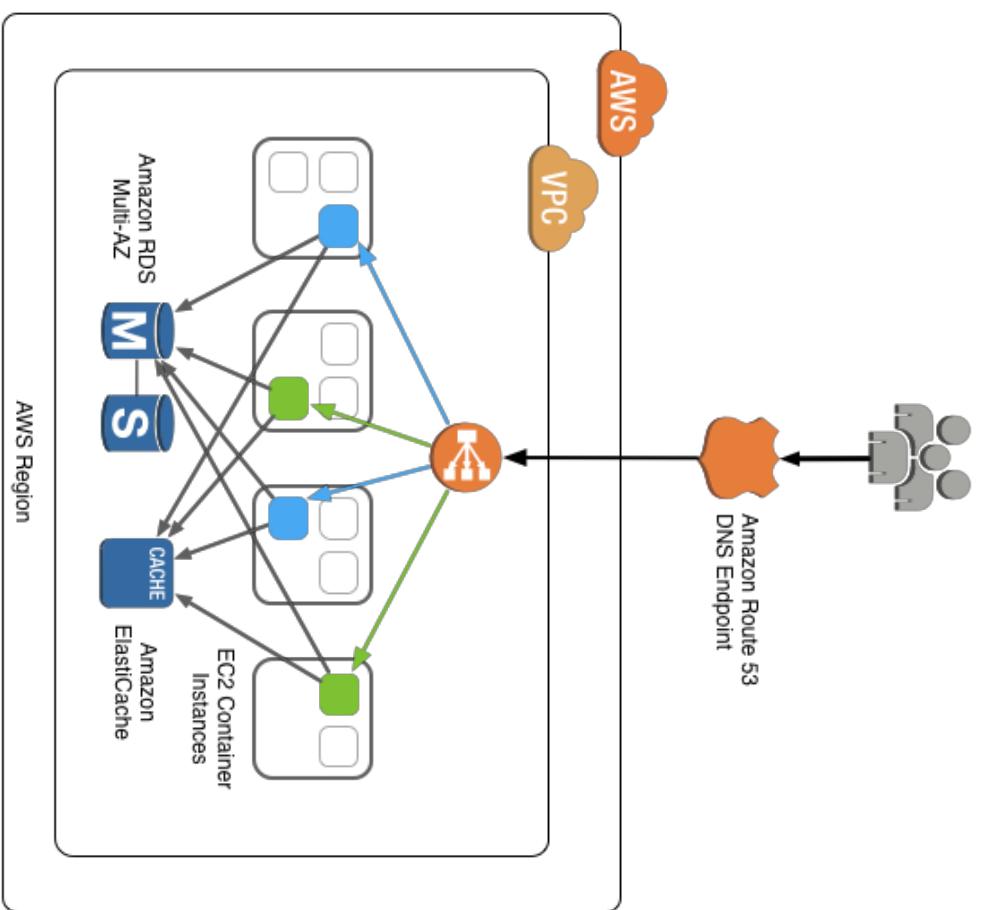
- Start with **blue** service composed of a task definition and ELB
- Create new task definition based on new version of Docker image
- Create **green** service with new task definition and map to existing ELB
- Scale up green service by incrementing number of tasks
- Decommission blue service by setting task count to 0



# Pattern: Swap ECS services with ELB

## Deployment process:

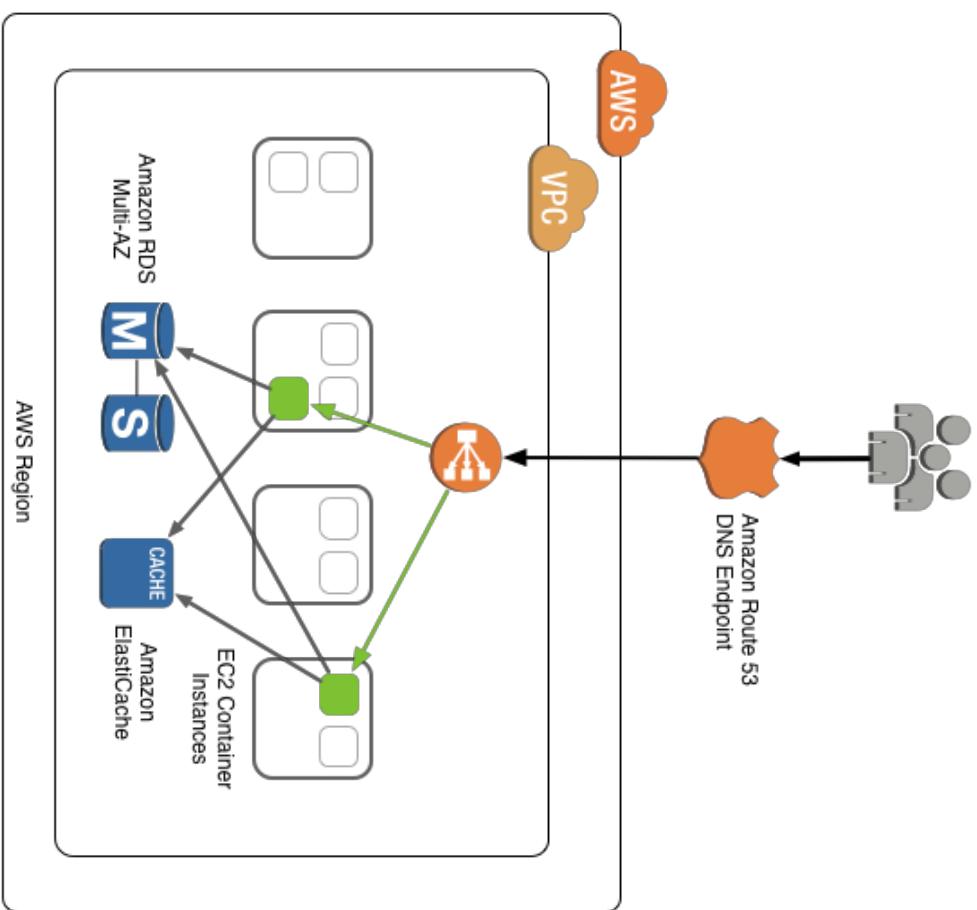
- Start with **blue** service composed of a task definition and ELB
- Create new task definition based on new version of Docker image
- Create **green** service with new task definition and map to existing ELB
- Scale up green service by incrementing number of tasks
- Decommission blue service by setting task count to 0



# Pattern: Swap ECS services with ELB

## Deployment process:

- Start with **blue** service composed of a task definition and ELB
- Create new task definition based on new version of Docker image
- Create **green** service with new task definition and map to existing ELB
- Scale up green service by incrementing number of tasks
- Decommission blue service by setting task count to 0



# Swapping ECS services with ELB

```
$ aws ecs create-service \
--cluster dice-demo \
--service-name web-service-v2 \
--task-definition web-tier \
--load-balancers loadBalancerName=web-
tier,containerName=web-tier,containerPort=8000 \
--desired-count 1 \
--role "ecsserviceRole"
```

# Swapping ECS services with ELB

```
$ aws ecs update-service \
--cluster dice-demo \
--service web-service-v2 \
--desired-count 2

$ aws ecs update-service \
--cluster dice-demo \
--service web-service-v1 \
--desired-count 0
```

- Similar to swapping ASG behind ELB
- Application versions must co-exist
- Container resources may need to be cleaned up

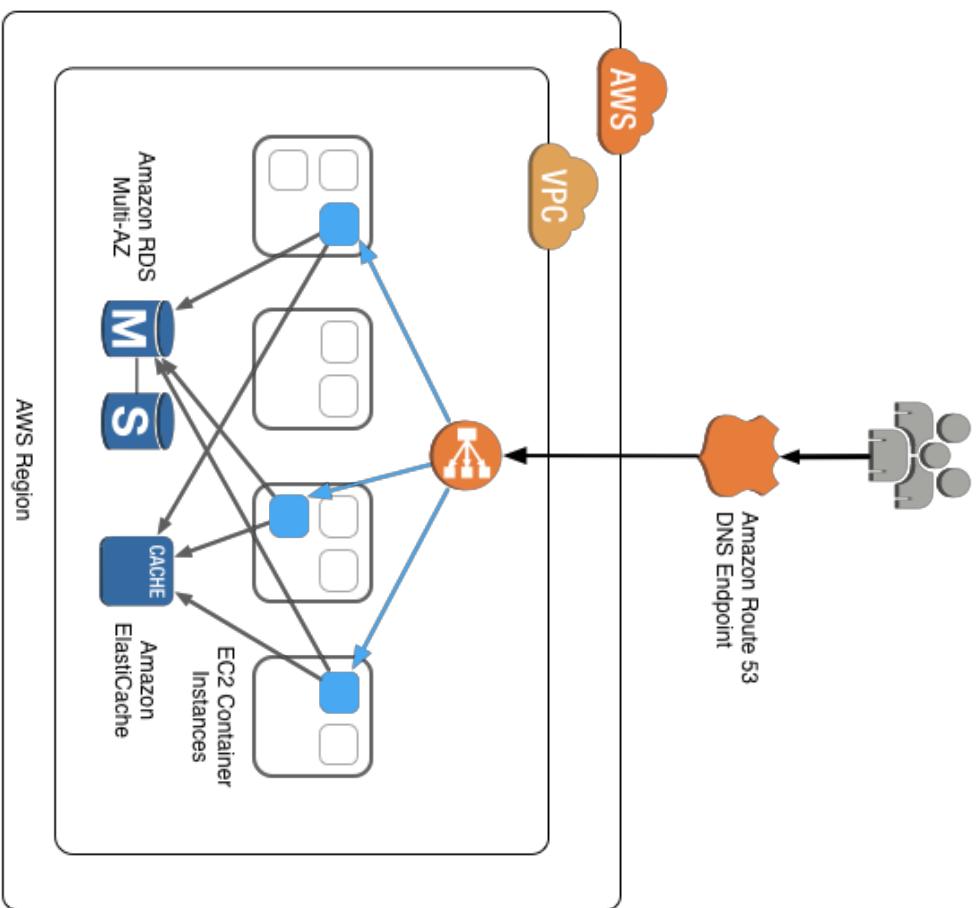
# Pattern review: Swap ECS services with ELB

| Risk Category           | Mitigation Level | Reasoning                                                     |
|-------------------------|------------------|---------------------------------------------------------------|
| Application issues      | Great            | Canary analysis                                               |
| Application performance | Great            | ELB already warm, CloudWatch and automation to scale ECS      |
| People/process errors   | Good             | Multi-step process to transition traffic between environments |
| Infrastructure failure  | Good             | Leverage CloudWatch, Auto Scaling, ELB                        |
| Rollback                | Great            | No DNS complexities                                           |
| Cost                    | Good             | Resource management handled by ECS                            |

# Pattern: ECS service update

## Deployment process:

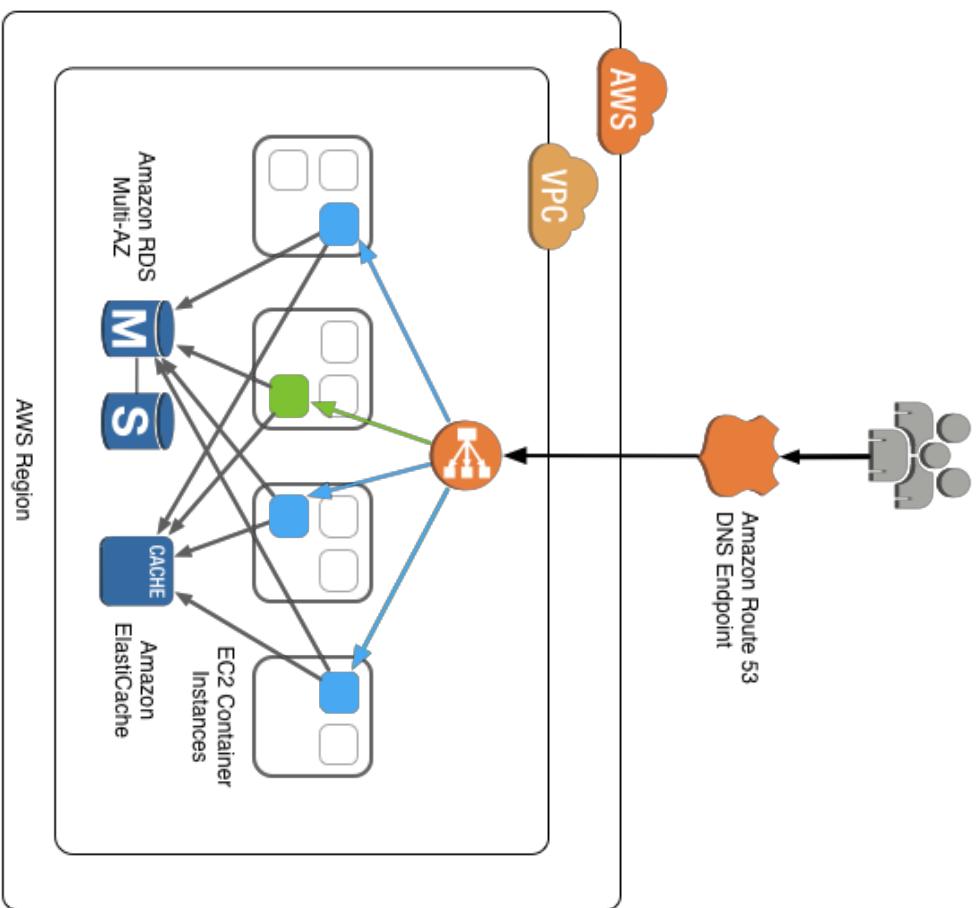
- Start with **blue** task definition referenced by an ECS service
- Create a **green** revision of the existing task definition
- Update existing ECS service to use the updated task definition
- ECS will deploy the new task definition to container instances in a rolling fashion



# Pattern: ECS service update

## Deployment process:

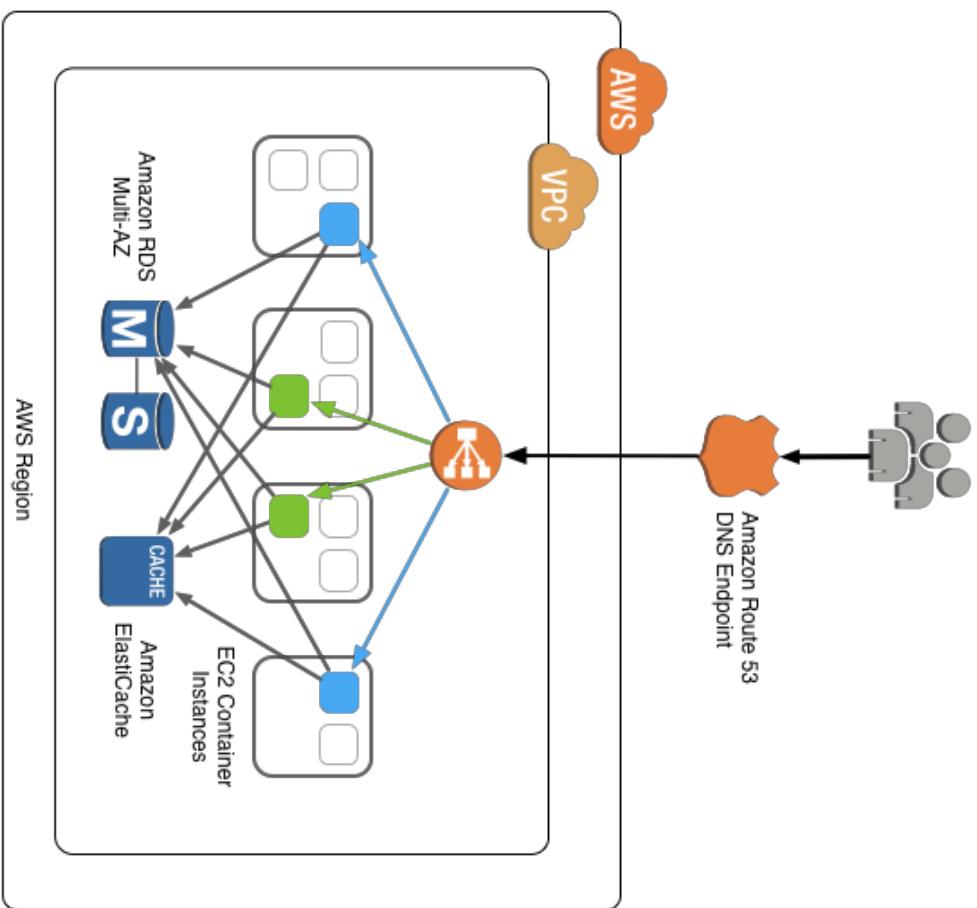
- Start with **blue** task definition referenced by an ECS service
- Create a **green** revision of the existing task definition
- Update existing ECS service to use the updated task definition
- ECS will deploy the new task definition to container instances in a rolling fashion



# Pattern: ECS service update

## Deployment process:

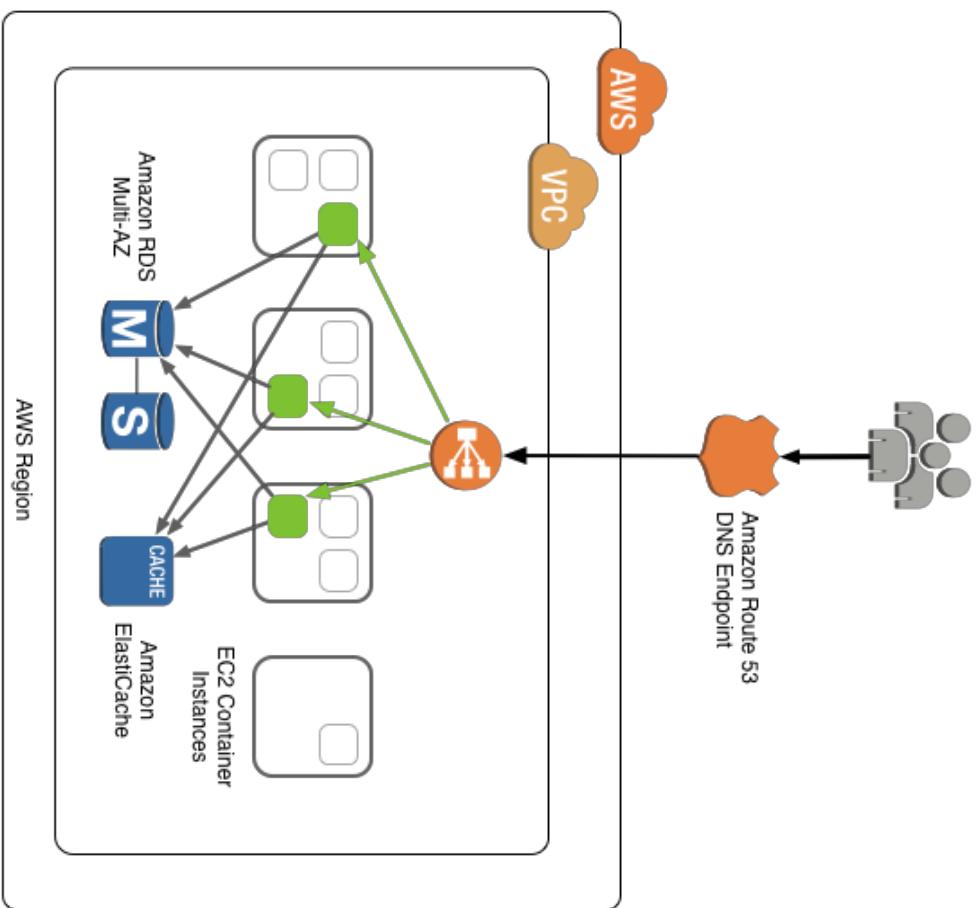
- Start with **blue** task definition referenced by an ECS service
- Create a **green** revision of the existing task definition
- Update existing ECS service to use the updated task definition
- ECS will deploy the new task definition to container instances in a rolling fashion



# Pattern: ECS service update

## Deployment process:

- Start with **blue** task definition referenced by an ECS service
- Create a **green** revision of the existing task definition
- Update existing ECS service to use the updated task definition
- ECS will deploy the new task definition to container instances in a rolling fashion



# ECS service update

## 1. Create task definition and update service

```
$ aws ecs update-service --cluster dice-demo --service web-service --task-definition
web-tier
```

## 2. Monitor service update process

```
$ aws ecs describe-services --cluster dice-demo --services web-service
```

```
{
 "services": [
 {
 "status": "ACTIVE",
 "taskDefinitiontier:2",
 }
]
}
```

# ECS service update

“deployments” section of describe-services output displays progress

```
"deployments": [
 {
 "status": "PRIMARY",
 "desiredCount": 3,
 "taskDefinition": "arn:aws:ecs:us-west-2:012345678901:task-definition/web-tier:2",
 "runningCount": 0
 },
 {
 "status": "ACTIVE",
 "desiredCount": 3,
 "taskDefinition": "arn:aws:ecs:us-west-2:012345678901:task-definition/web-tier:1",
 "runningCount": 3
 }
]
```

# ECS service update

“events” section of describe-services output will reveal any errors

```
"events": [
 {
 "message": "(service web-service) has reached a steady state.",
 "id": "34d17afa-89eb-454a-9311-f4de40222ca",
 "createdAt": 1442976067.928
 },
 {
 "message": "(service web-service) registered 1 instances in (elb web-tier)",
 "id": "b0020deb-8c4c-4128-bbd3-9b87487eba8f",
 "createdAt": 1442976061.644
 },
]
```

# Pattern: ECS service update

| Risk category           | Mitigation level | Reasoning                                 |
|-------------------------|------------------|-------------------------------------------|
| Application Issues      | Fair             | No canary analysis                        |
| Application Performance | Good             | No traffic management, ELB already warm   |
| People/Process Errors   | Great            | Simple, automated process                 |
| Infrastructure Failure  | Good             | Autoscaling cluster instances and service |
| Rollback                | Great            | No DNS complexities                       |
| Cost                    | Great            | Rolling deployment, cluster instance +1   |



What about schema changes?

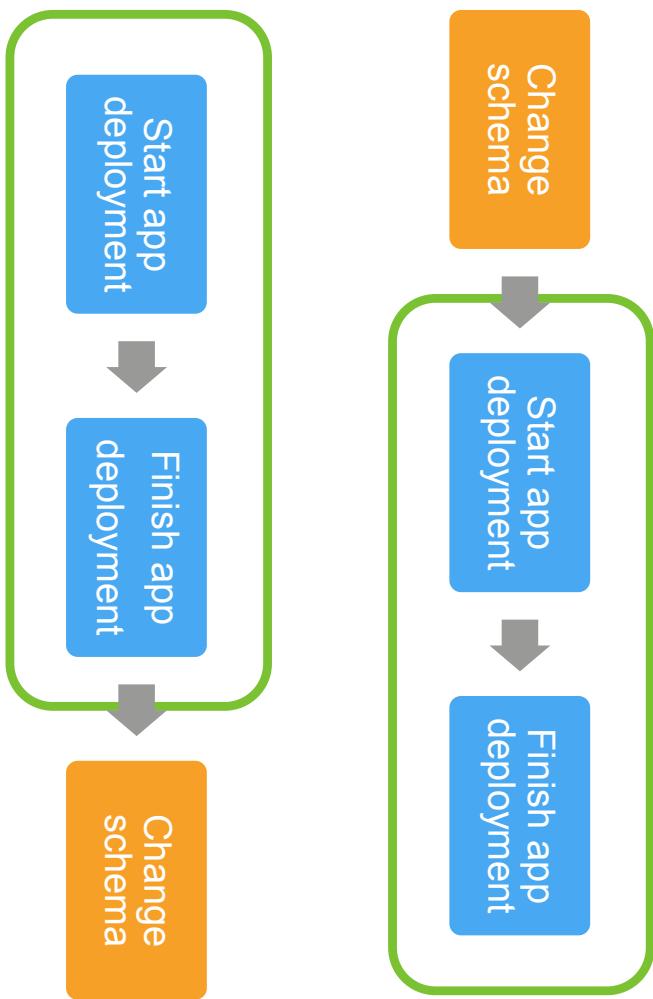
# Decoupled schema changes & code changes

## Two approaches:

- Database updates are backward-compatible (old code uses new schema)
- Code changes are backward-compatible with the old schema (new code uses old schema)

DB **outside** environment boundary

Tradeoff: **simplicity** vs. **risk**

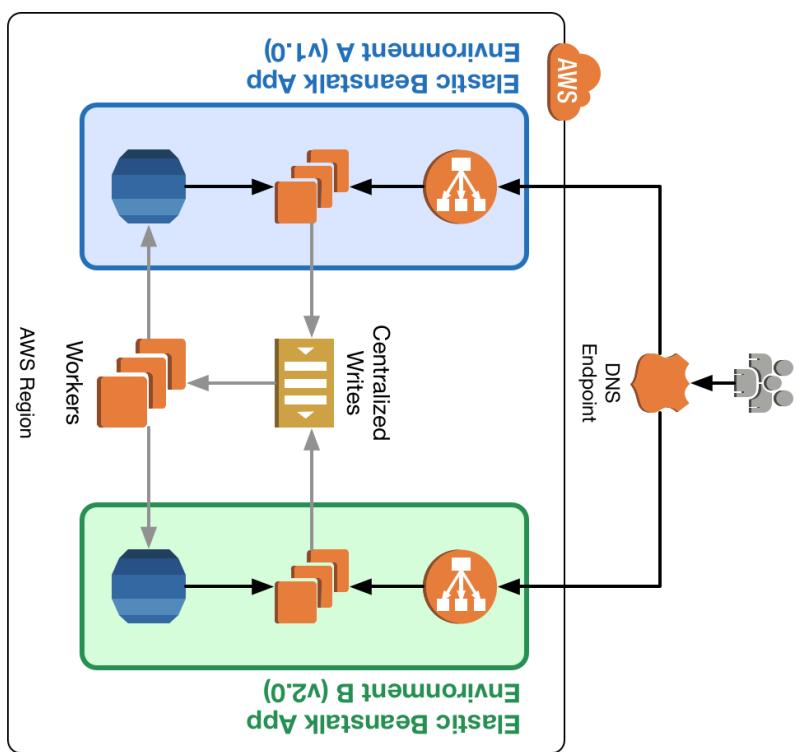




What if schema changes can't  
be decoupled, or you're  
deploying across regions?

# Isolated and synchronized data stores

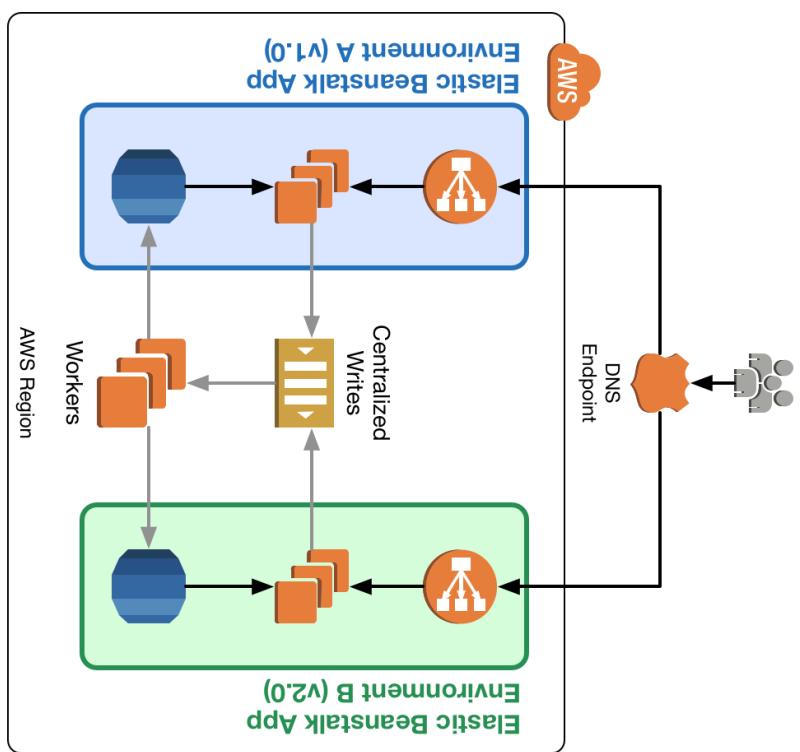
- DB **inside** environment boundary
- Environments have separate data stores
- Need to coordinate data changes across stacks:
  - Blue needs data for rollback
  - Green is new production stack
- What consistency model does the application use?



# Centralized writes process

## Deployment pattern:

- Deploy an aggregator of **centralized write operations**
- Common baseline of data: Use data store-specific replication
- Aggregator worker affects changes to data in both environments
- Account for lag/latency considerations



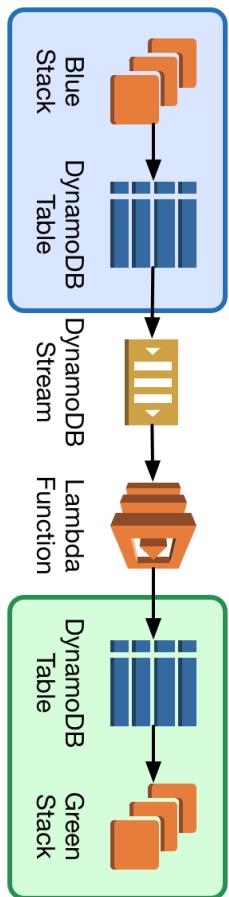
# Simplifying the centralized writes process

## Green app writes to both DBs

- Asynchronous process moves changes by blue app to green DB
- Great if there's a mismatch in consistency requirements (old app needs strong consistency)

## Each app writes to its own DB

- Asynchronous process pushes changes to the other DB
- Fully decoupled architecture
- Example: DynamoDB tables + streams + triggers + Lambda functions





# Closing thoughts

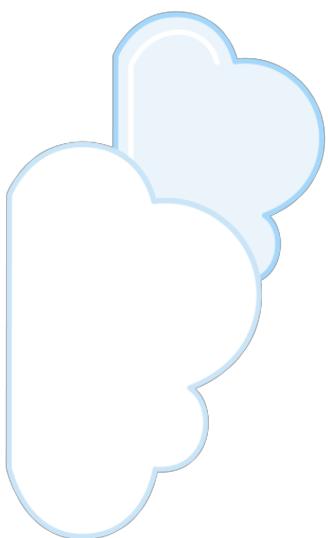
# Blue/green deployment patterns at a glance

| Pattern                 | Classic DNS cutover                                                         | Swap Auto Scaling groups   | Swap launch             | Swap ECS services with DNS     | Swap ECS services with ELB            | ECS service update              |
|-------------------------|-----------------------------------------------------------------------------|----------------------------|-------------------------|--------------------------------|---------------------------------------|---------------------------------|
| Application issues      | Canary analysis                                                             | Canary analysis            | Mixed fleet             | Canary analysis                | Canary analysis                       | Mixed fleet                     |
| Application performance | Granular traffic switch                                                     | Instance-level granularity | Mixed fleet             | ELB may require pre-warm       | Container level granularity, warm ELB | No traffic management, warm ELB |
| People/process errors   | Automation: Use CloudFormation with Elastic Beanstalk, OpsWork, third party |                            | Simple process DNS swap | Multi-step process             | ECS automated                         |                                 |
| Infrastructure failure  | Automation framework                                                        | Auto Scaling, ELB          | Auto Scaling, ELB       | CloudWatch, Auto Scaling, ELB  | CloudWatch, Auto Scaling, ELB         | CloudWatch, Auto Scaling, ELB   |
| Rollback capability     | DNS                                                                         | ELB                        | ELB                     | DNS                            | ECS automated                         | ECS automated                   |
| Cost management         | Gradual scaling                                                             | Gradual scaling            | Some over-provisioning  | Require addl cluster instances | Resource reuse                        | Rolling deployment              |
| Deployment complexity   | Simple, DNS weights                                                         | Auto Scaling control       | Scale-in adjustments    | Simple, DNS weights            | Multi-step process                    | Highly automated                |

# Get comfortable with the deployment process

**Deployments will always have risks associated with them**

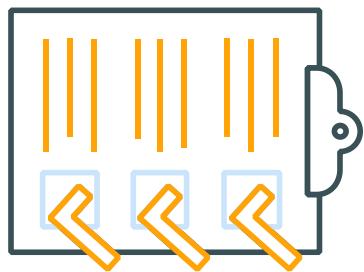
- Deployment & automation frameworks help mitigate process & human error risks
- Get comfortable with deployments, practice using production-like replicas
- AWS provides affordable, quickly provisioned resources – use the flexibility to experiment new approaches



Thank you!

AWS re:Invent

**Remember to complete  
your evaluations!**



## Related Sessions

DVO202 - DevOps at Amazon: A Look at Our Tools and Processes

DVO305 - Turbocharge Your Continuous Deployment Pipeline with Containers